
Decoupling What from Where: A Mechanism Study of Action-Type Supervision for GUI Grounding in a Small Vision-Language Model

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A GUI agent decides which action to take and where on the screen to take it; native
2 agent models emit both in one token stream. We ask what supervising the
3 action type does for the grounding half of that decision, using Qwen2-VL-2B
4 with LoRA on Android in the Wild and Multimodal-Mind2Web. A flat baseline
5 is compared with five ways of supplying the type under matched data, compute,
6 adapters, and decoding: an auxiliary loss, a hard-routed action word, an additive
7 learned embedding, a prepended learned token, and the type written into the
8 prompt. With five seeds, a cluster bootstrap over validation examples, and seed-
9 level paired tests, two results hold. On a training stream that mixes clicks and
10 scrolls with type events whose recorded target is a fixed off-screen point, the aux-
11 iliary loss, the additive embedding, and the type written into the prompt each raise
12 click grounding by 6 to 10 points over the baseline; the prepended learned token,
13 the mechanism we originally hypothesized, is no better than the baseline. Part of
14 the benefit is protection from the degenerate class: removing those events from
15 training raises the baseline itself by four points, after which the auxiliary loss’s
16 advantage is within noise and the additive embedding keeps a smaller one, and
17 on a separate stream of taps and swipes no mechanism helps. Interventions on
18 the trained models show that both learned embeddings are consulted at inference,
19 with a difference that matters for deployment: with its embedding zeroed, the
20 additive model returns to the baseline’s accuracy, whereas the prepended-token
21 model falls below it. We also correct our own description of a silent failure: in-
22 jecting the conditioning through `inputs_embeds` makes Qwen2-VL fall back to
23 one-dimensional positions for image tokens, which produced a nine-point deficit
24 that vanished once `input_ids` were preserved.

1 Introduction

26 A GUI agent maps a screenshot and an instruction to a low-level action: an action type such as *click*,
27 *scroll*, or *type*, and, for spatial actions, a target location. Current native agent models (UI-TARS [Qin
28 et al., 2025], OS-Atlas [Wu et al., 2025c], SeeClick [Cheng et al., 2024], ShowUI [Lin et al., 2025])
29 produce both in a single autoregressive stream. The two decisions are different kinds of prediction:
30 the action type is a small classification problem that depends on the goal and the screen state, and
31 the location is a localization problem in which the model has to find the evidence on the screen that
32 supports the action and commit to a point. Several recent agents separate the two, predicting the
33 type first and the target conditioned on it [Ma et al., 2024, Christianos et al., 2025]. What has not
34 been measured is what the type supervision itself does to the grounding model, and which of the
35 many ways of delivering it matters.

This paper is a matched-compute study of that question in a small model. We fine-tune Qwen2-VL-2B [Wang et al., 2024] with LoRA [Hu et al., 2022] to emit a coordinate string, and compare a flat baseline with five conditioning mechanisms that receive the action type in different forms: as an auxiliary classification loss, as a hard-routed action word in the output, as an additive learned embedding, as a prepended learned token, and as a word in the prompt. Each variant sees the same data in the same order and uses the same optimizer, adapters, and decoding. We started from the hypothesis that the prepended token would be the right mechanism, because it gives the model a dedicated, attendable position for the action type.

The results support type supervision and refute our hypothesis about the mechanism, and they depend on a property of the data that we did not appreciate at first. Android in the Wild records type events with a touch point off the screen, which our coordinate serializer clamps to the origin, so a training stream that mixes clicks, scrolls, and type events contains a class whose target is a fixed point. On that stream, the auxiliary loss and the additive embedding raise hit@0.10 by +0.064 and +0.054 over the flat baseline (five seeds, cluster bootstrap over examples, seed-level $p = 0.004$ and 0.028), almost entirely on clicks; the type written into the prompt as a word, with no new parameters, does as well, and the prepended token is within noise of the baseline. The obvious explanation, that the baseline learns to emit the off-screen point for clicks, accounts for little of the gap (it does so on 4% of clicks), but the class does hurt the baseline more broadly: removing type events from training while scoring the identical validation examples raises the baseline by four points, after which the auxiliary loss’s advantage is within noise and the additive embedding keeps a smaller one, and on a separate stream of taps and swipes neither helps.

Interventions on the trained models explain the ranking. A wrong type collapses either embedding model and a zeroed or class-averaged embedding costs accuracy in both, so both learned embeddings are consulted at inference; the difference is what is left when the signal is removed. The zeroed additive model sits at the flat baseline’s level, the zeroed prepended-token model seven points below it. An earlier implementation that injected the conditioning through `inputs_embeds`, which makes Qwen2-VL fall back to one-dimensional positions for its image tokens, scored nine points below the baseline, a result we initially read as evidence against conditioning.

Our contributions are (1) a matched-compute comparison of five action-type conditioning mechanisms against a flat baseline, with five seeds, a cluster bootstrap over examples, seed-level tests, and a control stream where conditioning should not help; (2) a diagnosis of where the gain comes from, including retraining on the same stream with the degenerate class removed; (3) interventions and attention measurements that separate mechanisms that use the type at inference from those that only use it in training; and (4) a corrected account of a positional-encoding failure that inverted a conclusion in a small model.

2 Related Work

GUI grounding and agent models. Mind2Web [Deng et al., 2023] and Android in the Wild [Rawles et al., 2023] supply the web and mobile episodes we train on; AndroidControl [Li et al., 2024] and AndroidWorld [Rawles et al., 2025] extend mobile evaluation. SeeClick [Cheng et al., 2024] argued that grounding is the bottleneck for visual GUI agents and introduced ScreenSpot, which ScreenSpot-Pro [Li et al., 2025] extends. UGround [Gou et al., 2025] and OS-Atlas [Wu et al., 2025c] scale grounding pretraining; Aguviz [Xu et al., 2025], UI-TARS [Qin et al., 2025], CogAgent [Hong et al., 2024], ScreenAI [Baechler et al., 2024], and Ferret-UI [You et al., 2024b] emit action and location together, the design our flat baseline follows. GUI-Actor [Wu et al., 2025a] replaces coordinate text with an attention head, UI-R1 [Lu et al., 2026] rewards type and argument separately, RegionFocus [Luo et al., 2025] zooms at test time, and OmniParser [Lu et al., 2024] and Set-of-Mark [Yang et al., 2023] externalize grounding into a parser.

Type-first agents. CoCo-Agent [Ma et al., 2024] decomposes AITW action prediction into the type followed by the target conditioned on it, in one stream; that is our hard-routing variant. Li-MAC [Christianos et al., 2025] places a small action-type transformer ahead of a LoRA-tuned Qwen2-VL-2B on AITW and AndroidControl and selects click targets conditioned on the predicted type, using screen elements rather than pixels. SeeAct [Zheng et al., 2024], Ponder & Press [Wang et al., 2025], and Aria-UI [Yang et al., 2025] separate text planning from pixel grounding; GUI-Libra [Yang et al., 2026] reweights reasoning against action tokens and GUI-AIMA [Zhou et al.,




Figure 1: Two-stage pipeline. Stage 1 classifies the action type from mean-pooled frozen features. Stage 2 grounds the action to a coordinate string, and the action type enters it through one of the mechanisms in the bottom row (D-text, the type as a word in the prompt, is not drawn). All mechanisms share data, compute, adapters, and decoding. D-slot is the retracted variant whose `inputs_embeds` injection bypassed the position computation.

2025] adds an anchor token for grounding. These systems establish that type-first pipelines work; none holds data, compute, and decoding fixed across conditioning mechanisms or intervenes on the trained model to ask whether the type signal is used, which is what this paper does.

Conditioning, faithfulness, and the backbone. SayCan [Ahn et al., 2022], ReAct [Yao et al., 2023], and RT-2 [Brohan et al., 2023] are the reference points for factored, exposed, and token-level action decisions; AutoUI [Zhang and Zhang, 2024], AppAgent [Zhang et al., 2025], DigiRL [Bai et al., 2024], WebRL [Qi et al., 2025], and VSC-RL [Wu et al., 2025b] improve mobile agents by imitation or reinforcement learning without changing how the type enters the grounding computation. Our auxiliary-loss variant is a multi-task objective in the sense of Caruana [1997]. Kosmos-2 [Peng et al., 2024], Shikra [Chen et al., 2023], and Ferret [You et al., 2024a] train VLMs to refer to image regions, and POPE [Li et al., 2023] and HallusionBench [Guan et al., 2024] ask whether an output is supported by the image; our attention measurements ask that of a coordinate. Qwen2-VL [Wang et al., 2024] assigns three-dimensional rotary positions (extending RoPE [Su et al., 2024]) to visual tokens from `input_ids`; Section 6.2 describes what happens when conditioning bypasses that path.

3 Method

3.1 Action taxonomy and grounding target

Both datasets are mapped onto a canonical eight-class action space (*click*, *double-click*, *type*, *scroll*, *drag*, *hotkey*, *wait*, *finished*). Mind2Web’s CLICK, SELECT, and HOVER map to *click* and TYPE to *type*. AITW stores gestures as touch and lift points; we split them into taps (*click*) and swipes (*scroll*) by the normalized distance between the two, with threshold 0.04. The grounding target is the touch point normalized to $[0, 1]^2$, serialized as (x, y) with integers on a 0–999 grid so every target has the same token length. AITW records non-touch actions such as typing with the touch point $(-1, -1)$; the serializer clamps this to $(0, 0)$, so every *type* event in the training stream has the same target at the screen origin, and its distance to the raw target is exactly $\sqrt{2}$ whatever the model predicts. This detail matters for the analysis in Section 5.1.

3.2 Stage 1: which action

Stage 1 runs the frozen Qwen2-VL-2B-Instruct backbone over the screenshot and instruction, mean-pools the final hidden states over unmasked tokens into a 1536-dimensional vector, and trains a three-layer MLP ($1536 \rightarrow 1024 \rightarrow 256 \rightarrow 8$, dropout 0.1) with class-weighted cross-entropy, AdamW (learning rate 10^{-4} , weight decay 0.01), batch size 128, for eight epochs, keeping the checkpoint with the best validation macro-F1. We compare features from text only, text plus screenshot, and a sanity variant with zeroed features.

122 3.3 Stage 2: where

123 Stage 2 fine-tunes the same backbone with LoRA adapters (rank 16, $\alpha = 32$, dropout 0.05) on the
 124 query, key, value, and output projections, 4.4M trainable parameters out of 2.2B. The prompt is the
 125 screenshot, `Goal: {goal}`, and a fixed instruction to predict the action coordinate; the answer is
 126 the coordinate string. With s the screenshot, g the goal, a the action type, and $y_{1:T}$ the coordinate
 127 tokens, every variant minimizes

$$\mathcal{L}_{\text{coord}} = - \sum_{t=1}^T \log p_{\theta}(y_t | y_{<t}, s, g, c), \quad (1)$$

128 masked to the answer tokens, with c the variant-specific conditioning.

129 **A, flat.** $c = \emptyset$: the joint-decoding design of native agent models, reduced to the grounding sub-
 130 problem.

131 **B, auxiliary loss.** A linear head h_{ϕ} on the mean-pooled last hidden state $\bar{z}$ predicts the action type
 132 during training, $\mathcal{L}_B = \mathcal{L}_{\text{coord}} + \lambda \text{CE}(h_{\phi}(\bar{z}), a)$ with $\lambda = 1$. Nothing about the type is provided at
 133 inference, so any gain is a training effect.

134 **C, hard routing.** The model is trained to emit the action word before the coordinate, `click (x,`
 135 `y)`, as in CoCo-Agent’s conditional action prediction. At test time the word for the gold (or Stage-1)
 136 type is forced as a decoding prefix.

137 **D-hook, additive embedding.** A learned table $E \in \mathbb{R}^{8 \times d}$ holds one vector per class; a forward hook
 138 on the embedding layer adds $E[a]$ to every text and image token embedding. E is zero-initialized,
 139 so D-hook starts as an exact copy of A.

140 **D-token, prepended action token.** A new token `<|action_slot|>` is added to the vocabulary
 141 and placed at the start of the text prompt, so it occupies a real position in `input_ids`; a forward
 142 hook replaces its embedding with $E[a]$, initialized from $\mathcal{N}(0, 0.02^2)$. This is the architecture we
 143 originally hypothesized: a full-norm, attendable representation of the type.

144 **D-text, type as a word.** The prompt begins with `Action: click.` (or the gold type’s word),
 145 using the pretrained word embeddings and no new parameters. This is the obvious baseline for the
 146 two learned-embedding variants.

147 For C, D-hook, D-token, and D-text the gold type is used in training and at evaluation unless stated
 148 otherwise; Section 5.3 closes the loop with Stage-1 predictions. All variants share the data order, the
 149 optimizer (AdamW, learning rate 2×10^{-5} , weight decay 0.01, batch size 1, gradient clipping at 1.0,
 150 two epochs), and greedy decoding with 16 new tokens. No hyperparameter was tuned per variant;
 151 Section 7 returns to this.

152 4 Experimental setup

153 **Data.** Stage 1 uses 7,775 Multimodal-Mind2Web training steps evaluated on the 1,338-step
 154 `test_task` split, and 5,000 AITW steps from the General subset evaluated on 1,000 held-out
 155 steps. Stage 2 uses two AITW streams built by filtering the training split in its stored order.
 156 `all_with_coords` keeps taps, swipes, and type events; at 1,200 training examples its validation
 157 slice of 250 steps holds 169 clicks, 46 scrolls, and 35 type events. `taps_and_swipes` keeps only
 158 taps and swipes (1,000 training and 200 validation steps); both land anywhere on the screen. The
 159 AITW stream is deterministic, so for a given mix and training size every seed and every variant
 160 is evaluated on the same examples in the same order, which is what makes paired tests possible.
 161 The validation slice begins where the training slice ends; consecutive steps come from the same
 162 episodes, and absolute numbers are not comparable across training sizes. A second control uses
 163 Multimodal-Mind2Web (1,000 training and 250 evaluation steps), where the target is the center of
 164 the gold element. The Stage-2 study is run on AITW because that is where the action type is a vision
 165 problem: Stage 1 (Table 5 in the appendix) gains +0.414 macro-F1 from the screenshot on AITW,
 166 whose instructions are goals, and +0.009 on Mind2Web, whose instructions name the action.

167 **Metrics.** `hit@r` is the fraction of predictions within normalized Euclidean distance r of the target,
 168 for $r \in \{0.05, 0.10, 0.25\}$; `hit@0.10` is the headline. We also report the mean normalized distance.

Table 1: Stage 2 grounding on AITW all_with_coords ($n_{\text{train}}=1200$, $n_{\text{val}}=250$, 5 seeds). Mean $\pm$ std over seeds; Δ with 95% paired-bootstrap CI over pooled (seed, example) units; * $p < .05$, ** $p < .01$, *** $p < .001$.

Variant	hit@0.05	hit@0.10	hit@0.25	mean L2 $\downarrow$	Δ hit@0.10 vs. A [cluster CI]	seed-level p
A: flat	0.130 $\pm$ 0.039	0.229 $\pm$ 0.043	0.499 $\pm$ 0.029	0.406 $\pm$ 0.028		
B: aux. loss	0.178 $\pm$ 0.025	0.293 $\pm$ 0.021	0.582 $\pm$ 0.026	0.365 $\pm$ 0.004	+0.064 [+0.036, +0.093]***	0.004
C: hard routing	0.170 $\pm$ 0.022	0.264 $\pm$ 0.045	0.538 $\pm$ 0.059	0.421 $\pm$ 0.016	+0.035 [+0.001, +0.069]*	0.189
D-hook: additive	0.170 $\pm$ 0.042	0.282 $\pm$ 0.041	0.563 $\pm$ 0.016	0.370 $\pm$ 0.010	+0.054 [+0.022, +0.086]***	0.028
D-token: prepended	0.144 $\pm$ 0.056	0.238 $\pm$ 0.049	0.504 $\pm$ 0.055	0.386 $\pm$ 0.021	+0.009 [-0.015, +0.034]	0.465
D-text: word in prompt	0.184 $\pm$ 0.024	0.304 $\pm$ 0.032	0.562 $\pm$ 0.035	0.365 $\pm$ 0.009	+0.075 [+0.034, +0.117]***	0.052

Unparseable outputs count as misses at distance $\sqrt{2}$; only hard routing produces them (2 to 8% of outputs), and its numbers are reported both ways. Per-class numbers are computed on the same basis as the overall ones.

Statistics. The headline setting uses five seeds (42 to 46); every other cell uses three (42 to 44). We report the mean and sample standard deviation over seeds. Every run logs a per-example distance on a shared evaluation set, and every seed of every variant shares the per-seed data order, so comparisons are paired at two levels. Our primary interval is a cluster bootstrap that resamples validation examples with replacement and keeps all seeds of a sampled example together (10,000 resamples, 95% percentile interval, two-sided bootstrap p [Efron and Tibshirani, 1993]); resampling (seed, example) units separately would treat the seeds of one example as independent and gives intervals about a third narrower. We also report a paired t -test on the five (or three) per-seed means, which has little power but cannot be fooled by within-seed correlation, and treat a difference as established when both agree; stars come from the cluster bootstrap with the seed-level p beside them. No correction is made for the number of contrasts; the scaling study in Section 5.4 is presented as descriptive for that reason.

5 Results

5.1 Type supervision improves click grounding on the mixed stream

Table 1 is the main result: six variants on all_with_coords at 1,200 training examples, five seeds each. The auxiliary loss (B) and the additive embedding (D-hook) beat the flat baseline on every metric, by +0.064 and +0.054 hit@0.10 with cluster intervals that exclude zero and seed-level p of 0.004 and 0.028, and they are indistinguishable from each other (D-hook – B: -0.010 , cluster CI $[-0.043, +0.024]$, seed $p = 0.56$). Hard routing (C) gains +0.035 with a cluster interval that barely excludes zero and a seed-level p of 0.19, and loses on mean distance, partly because forcing the action word makes 2 to 8% of its outputs unparseable; scored on parsed outputs only, its click rate is unchanged (0.356 vs. 0.350). The prepended token (D-token), our hypothesized mechanism, is within noise of the baseline on every hit rate (+0.009, cluster CI $[-0.015, +0.034]$) and below B by -0.055 (cluster CI $[-0.088, -0.024]$, seed $p = 0.016$). The type written into the prompt as a word (D-text), which adds no parameters, gains +0.057 over the baseline (cluster CI $[+0.031, +0.084]$, seed $p = 0.04$) and is within noise of B (+0.007, cluster CI $[-0.020, +0.035]$), so the failure is specific to the learned prepended token, not to conditioning at the input. Seed variance is large: the baseline ranges from 0.169 to 0.280 across seeds, which is why we report five seeds and paired statistics rather than a best run. The ordering with the first three seeds was the same; the two added seeds lowered every variant and widened the gap between D-token and the two working mechanisms.

Table 2 shows where the gain lives and tests the simplest explanation for it. The improvement is a click improvement: B and C gain ten points on clicks, D-hook six, and C pays for its click gain with a sixteen-point loss on scrolls because forcing the word *scroll* into the output derails the gesture; D-hook is the only variant that improves both classes. Since the only difference between this stream and the control is the type class with its fixed target, one might expect the flat baseline’s click errors to be emissions of that target that the conditioned models avoid. They mostly are not. The baseline emits (0, 0) for 4.3% of clicks (13% on its worst seed), the auxiliary-loss model for 2.4%, and the three conditioned-at-inference models never; excluding those predictions moves the baseline’s click rate from 0.256 to 0.266 while B stays at 0.366 and D-hook at 0.321. The class degrades the baseline’s click grounding more broadly than by exact emissions, and Table 7 in the appendix

Table 2: Where the gain lives. Per-class hit@0.10 on the headline mix (five seeds, same basis as Table 1), the fraction of click predictions that are the clamped type target (0, 0), and click hit@0.10 after excluding those predictions. Every model scores zero on *type* because the raw target is off-screen.

Variant	click	scroll	type	clicks sent to (0, 0)	click, excluding those
A: flat	0.256 $\pm$ 0.062	0.304 $\pm$ 0.034	0.000	4.3% (13.0% on the worst seed)	0.266
B: aux. loss	0.357 $\pm$ 0.026	0.278 $\pm$ 0.018	0.000	2.4%	0.366
C: hard routing	0.350 $\pm$ 0.061	0.148 $\pm$ 0.073	0.000	0.0%	0.350
D-hook: additive	0.321 $\pm$ 0.064	0.357 $\pm$ 0.039	0.000	0.0%	0.321
D-token: prepended	0.269 $\pm$ 0.069	0.304 $\pm$ 0.090	0.000	0.0%	0.269

Table 3: Control: AITW `taps_and_swipes` ($n_{\text{train}}=1000$, $n_{\text{val}}=200$, 3 seeds), where action type is not spatially informative.

Variant	hit@0.05	hit@0.10	hit@0.25	mean L2 $\downarrow$	$\Delta\text{hit@0.10 vs. A}$ [cluster CI]	seed-level p
A: flat	0.243 $\pm$ 0.020	0.382 $\pm$ 0.015	0.720 $\pm$ 0.017	0.185 $\pm$ 0.003		
B: aux. loss	0.225 $\pm$ 0.015	0.368 $\pm$ 0.012	0.730 $\pm$ 0.039	0.173 $\pm$ 0.011	-0.013 [-0.038, +0.012]	0.456
C: hard routing	0.217 $\pm$ 0.016	0.325 $\pm$ 0.040	0.635 $\pm$ 0.064	0.271 $\pm$ 0.068	-0.057 [-0.108, -0.007]*	0.215
D-hook: additive	0.235 $\pm$ 0.022	0.363 $\pm$ 0.021	0.725 $\pm$ 0.043	0.182 $\pm$ 0.009	-0.018 [-0.052, +0.015]	0.334

measures how much: we retrain A, B, and D-hook on the same training slice with the type events removed and score the identical 250 validation examples (three seeds). Removing the class raises the flat baseline from 0.255 to 0.297 hit@0.10 (cluster CI on the change [+0.016, +0.069]; clicks from 0.294 to 0.365). On that type-free stream the auxiliary loss’s advantage over the stronger baseline is within noise (+0.021, cluster CI [-0.004, +0.045]) and the additive embedding keeps a smaller established one (+0.046, cluster CI [+0.008, +0.084]). With the control in the next section, we read this as follows: much of what type supervision buys on the mixed stream is protection of click grounding from a class whose target is degenerate; a residual advantage for the additive embedding appears in one type-free setting and not the other, so we do not claim a general spatial prior.

5.2 Control: no gain on taps and swipes

On `taps_and_swipes` (Table 3), whose training stream has no degenerate class, B and D-hook are within noise of the baseline (-0.013 and -0.018 hit@0.10, cluster intervals [-0.038, +0.012] and [-0.052, +0.015]), so the data admit harms of up to four points as well as no effect, and C is worse by -0.057 (cluster CI [-0.108, -0.007], seed $p = 0.22$) with a large increase in mean distance. The Multimodal-Mind2Web control sits at floor (a tenth of predicted points fall inside the gold element for both A and D-hook, three seeds) and is uninformative in either direction; we report it in the appendix and do not lean on it.

5.3 End to end: predicted types instead of gold

The tables above condition on the gold action type. To close the loop we train D-hook, fit the Stage-1 MLP on frozen features of its training examples for a fixed 200 full-batch epochs with no checkpoint selection, and evaluate D-hook on the same validation examples with gold types and with Stage-1 predictions (three seeds; Stage 1 reaches 0.843 accuracy). The oracle pipeline scores 0.292 hit@0.10, the predicted-type pipeline 0.271, and the flat baseline on the same seeds 0.255. Only the oracle-to-predicted gap is established (+0.021, cluster CI [+0.003, +0.041]); the predicted pipeline’s margin over the baseline (+0.016, cluster CI [-0.013, +0.047], seed $p = 0.44$) is not, and we report it as descriptive. Classifier error removes more than half of the conditioning gain in this setting.

5.4 Scaling: a descriptive study

Figure 2 extends the comparison to $n \in \{300, 500, 800, 1200, 2500, 5000\}$ with three seeds at every size. The cluster intervals for D-hook - A exclude zero at five of the six sizes on hit@0.10 (all six on hit@0.25), and its estimate at 5,000 examples (+0.043) is the same size as at 1,200; B - A excludes zero at 300 and 1,200, is negative at 800, and is within noise elsewhere. With three seeds the seed-level tests reach $p < 0.05$ only at 300 and 5,000 for D-hook and at 300 and 1,200 for B, and a direct D-hook - B contrast is within noise at four of the six sizes. We read the study as consistent




Figure 2: Conditioning advantage over the flat baseline as a function of training-set size on `all_with_coords`, three seeds per cell, 95% cluster-bootstrap intervals over examples. Each size is scored on a different validation slice with a different class mix (Table 8), so only within-size differences are meaningful, and 24 contrasts are shown without correction.

Table 4: Interventions on the trained conditioned models: the same model is decoded with the gold action embedding, a wrong one (cyclic permutation; `click→type`), a wrong one that swaps only clicks and scrolls, the table zeroed, and every row replaced by the class mean (uninformative, in-distribution norm). Five seeds $\times$ 250 examples; Δ = gold – condition with 95% cluster-bootstrap CI over examples and the seed-level paired p . Flat A on the same examples is the reference row.

Model	Conditioning	hit@0.10	click hit@0.10	scroll hit@0.10	Δ gold – cond. (hit@0.10)	seed p
D-hook: additive	gold type	0.275 ± 0.040	0.310 ± 0.060	0.357 ± 0.052		
	wrong (cyclic)	0.054 ± 0.009	0.006 ± 0.000	0.270 ± 0.050	$+0.222$ [$+0.174, +0.270$]***	0.000
	wrong (click↔scroll)	0.099 ± 0.056	0.073 ± 0.087	0.270 ± 0.050	$+0.176$ [$+0.133, +0.221$]***	0.003
	zeroed	0.248 ± 0.032	0.288 ± 0.053	0.291 ± 0.050	$+0.027$ [$-0.006, +0.060$]	0.224
	class mean	0.206 ± 0.058	0.217 ± 0.099	0.322 ± 0.064	$+0.070$ [$+0.034, +0.106$]***	0.133
	flat A, same examples	0.229 ± 0.043	0.256	0.304	$+0.046$ [$+0.017, +0.077$]**	0.032
D-token: prepended	gold type	0.251 ± 0.057	0.284 ± 0.085	0.321 ± 0.113		
	wrong (cyclic)	0.080 ± 0.014	0.036 ± 0.022	0.304 ± 0.061	$+0.171$ [$+0.125, +0.218$]***	0.005
	wrong (click↔scroll)	0.140 ± 0.030	0.124 ± 0.040	0.304 ± 0.061	$+0.111$ [$+0.065, +0.158$]***	0.045
	zeroed	0.159 ± 0.083	0.154 ± 0.129	0.299 ± 0.033	$+0.092$ [$+0.055, +0.130$]***	0.164
	class mean	0.165 ± 0.049	0.163 ± 0.076	0.299 ± 0.027	$+0.086$ [$+0.050, +0.123$]***	0.113
	flat A, same examples	0.229 ± 0.043	0.256	0.304	$+0.007$ [$-0.019, +0.034$]	0.667

with the additive embedding’s advantage persisting with data and the auxiliary loss’s not, and no stronger than that; an earlier single-seed curve had suggested that all conditioning gains shrink with data, which three seeds do not support. Figure 4 in the appendix shows individual screens.

6 Mechanism

6.1 Interventions on both embedding mechanisms

D-token’s failure to help could mean that the model ignores the slot, and D-hook’s success could be a training effect that leaves a vector the model never reads. To test both on equal terms we retrain each at the headline setting (five seeds) and decode every trained model under five conditionings on the full validation slice: the gold type; a wrong type from a cyclic permutation of the classes present (`click` to `type`, `type` to `scroll`, `scroll` to `click`); a second wrong map that swaps only clicks and scrolls, so no groundable example is routed to the degenerate class; the table zeroed; and every row replaced by the mean of the rows of the classes present, which is uninformative but in-distribution. Table 4 reports the results with the flat baseline on the same examples; an earlier three-condition test on separately trained D-token models (Table 9, appendix) agrees. The learned rows of D-hook are small: each has norm 0.04 against a mean token-embedding norm of 0.58, whereas D-token’s rows have norm 0.78. Small as they are, the model reads them. A wrong type collapses D-hook from 0.275 to 0.054 under the cyclic map and to 0.099 under the click-scroll map (cluster intervals [$+0.174, +0.270$] and [$+0.133, +0.221$], seed $p < 0.001$ and 0.003), so the collapse is not an artifact of routing clicks to the degenerate class. The class-mean vector costs seven points ($+0.070$, cluster CI [$+0.034, +0.106$], seed $p = 0.13$). Zeroing costs about three, and that difference is not established ($+0.027$, cluster CI [$-0.006, +0.060$], seed $p = 0.22$): the zeroed model sits at the level

of the flat baseline (+0.019 over A, not established). The single-seed attention run in Section ??, on which an earlier version of this paper based the claim that the embedding could be removed without loss, is consistent with these intervals and was simply underpowered. The additive embedding therefore behaves as a switch at inference: with the right type the model gains six points over the baseline, without a signal it returns to the baseline’s level, and with a wrong or blended signal it collapses.

D-token differs in one respect that matters for deployment. Every perturbation collapses it, including the in-distribution class-mean condition (+0.101, cluster CI [+0.057, +0.147]), so its sensitivity is not an artifact of an out-of-distribution zero vector, and the zeroed model is worse than the flat baseline by −0.075 (cluster CI [−0.112, −0.039]), whereas the zeroed D-hook is not. Adaptation with a prepended token makes the model depend on the token; adaptation with an additive vector does not, even though the vector is used. In both mechanisms the conditioning acts almost entirely on clicks. [re-check D-token numbers after seeds 45 and 46 land] Attention measured at the token that predicts the y coordinate moves the same way as accuracy under these interventions on a single seed (Appendix G); we treat those numbers as a consistency check, not as evidence about faithfulness on their own.

6.2 A positional-encoding failure that inverted a conclusion

Our first implementation of the prepended token replaced the slot’s embedding by building `inputs_embeds` and passing that to the model instead of `input_ids`. On `taps_and_swipes` that variant scored 0.298 ± 0.026 hit@0.10 against 0.390 ± 0.010 for the flat baseline in the same set of runs, and a variant with the same slot but a frozen random embedding scored the same (0.290), which pointed at the injection path. Qwen2-VL derives its three-dimensional rotary positions from `input_ids`; given `inputs_embeds` alone, the transformers implementation falls back to one-dimensional positions replicated across the three axes, and the image tokens lose their two-dimensional layout (details, version, and a related library report [Hugging Face Transformers contributors, 2024] in Appendix F). Keeping `input_ids` intact and modifying embeddings through a forward hook removed the deficit. We record this because the failure is silent, because we initially read it as evidence against conditioning, and because an earlier version of this paper described its mechanism incorrectly as a shift of the visual positions.

7 Limitations

The study is small: one 2B backbone, LoRA adapters, validation slices of 200 to 250 steps, one learning rate and one adapter configuration for every variant, and one L4 per run. The absolute grounding numbers are far below the field’s leaderboards, per-seed variance is large, and the conclusions rest on paired statistics reported at two levels because neither is clean: the cluster bootstrap ignores that seeds share a data order, and the seed-level test has five or three units. The refutation of the prepended token is conditional on that untuned configuration; a stronger D-token could exist under a schedule that trains the fresh token faster, and we did not train a control that keeps the slot but freezes its embedding. The headline stream’s only spatially informative class has a fixed off-screen target, so the claim is about adaptation on a mixed stream that contains such a class, not about a general spatial prior. The end-to-end result is descriptive, the scaling study scores each size on a different validation slice with 24 uncorrected contrasts, the attention analysis uses one seed and one layer with a probe token chosen after the fact, and Mind2Web is at floor. We do not evaluate on ScreenSpot or on multi-step task success, where the compounding-error motivation of the introduction would have to be tested directly. Code, split indices, and prompts will be released.

8 Conclusion

Supervising the action type during adaptation of a small grounding model helps when the training stream mixes groundable actions with a class whose target is degenerate, and the help is concentrated on clicks; much of it is protection from that class, since removing the class raises the flat baseline itself, and on a stream of taps and swipes conditioning does nothing. The mechanism ranking is stable across five seeds and two levels of paired inference: an auxiliary loss, an additive embedding, and the type written into the prompt capture the gain, the additive embedding keeps an advantage

across training sizes from 300 to 5,000 examples, and the prepended learned token we set out to validate is consulted by the model and improves nothing over a baseline that never saw a type. Both learned embeddings are read at inference and both collapse under a wrong type; what distinguishes them is that a model adapted with the additive vector still grounds at the baseline’s level when the vector is removed, while a model adapted with the prepended token does not. For practitioners adapting small VLMs as GUI agents on mixed action streams, the usable lessons are to supervise the action type, to prefer conditioning that degrades to the baseline rather than below it when the type is missing or wrong, and to keep `input_ids` on the path whenever a conditioning scheme touches the embeddings of a model with multimodal positional encoding.

References

- Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, Daniel Ho, Jasmine Hsu, Julian Ibarz, Brian Ichter, Alex Irpan, Eric Jang, Rosario Jauregui Ruano, Kyle Jeffrey, Sally Jesmonth, Nikhil J Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng Kuang, Kuang-Huei Lee, Sergey Levine, Yao Lu, Linda Luu, Carolina Parada, Peter Pastor, Jornell Quiambao, Kanishka Rao, Jarek Rettinghouse, Diego Reyes, Pierre Sermanet, Nicolas Sievers, Clayton Tan, Alexander Toshev, Vincent Vanhoucke, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Mengyuan Yan, and Andy Zeng. Do as I can, not as I say: Grounding language in robotic affordances. In *Conference on Robot Learning (CoRL)*, 2022.
- Gilles Baechler, Srinivas Sunkara, Maria Wang, Fedir Zubach, Hassan Mansoor, Vincent Etter, Victor Cărbune, Jason Lin, Jindong Chen, and Abhanshu Sharma. ScreenAI: A vision-language model for UI and infographics understanding. In *Proceedings of the Thirty-Third International Joint Conference on Artificial Intelligence (IJCAI)*, 2024.
- Hao Bai, Yifei Zhou, Mert Cemri, Jiayi Pan, Alane Suhr, Sergey Levine, and Aviral Kumar. DigiRL: Training in-the-wild device-control agents with autonomous reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choromanski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn, Pete Florence, Chuyuan Fu, Montse Gonzalez Arenas, Keerthana Gopalakrishnan, Kehang Han, Karol Hausman, Alexander Herzog, Jasmine Hsu, Brian Ichter, Alex Irpan, Nikhil Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng Kuang, Isabel Leal, Lisa Lee, Tsang-Wei Edward Lee, Sergey Levine, Yao Lu, Henryk Michalewski, Igor Mordatch, Karl Pertsch, Kanishka Rao, Krista Reymann, Michael Ryoo, Grecia Salazar, Pannag Sanketi, Pierre Sermanet, Jaspier Singh, Anikait Singh, Radu Soricut, Huong Tran, Vincent Vanhoucke, Quan Vuong, Ayzaan Wahid, Stefan Welker, Paul Wohlhart, Jialin Wu, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Tianhe Yu, and Brianna Zitkovich. RT-2: Vision-language-action models transfer web knowledge to robotic control. In *Conference on Robot Learning (CoRL)*, 2023.
- Rich Caruana. Multitask learning. *Machine Learning*, 28(1):41–75, 1997. doi: 10.1023/A:1007379606734.
- Keqin Chen, Zhao Zhang, Weili Zeng, Richong Zhang, Feng Zhu, and Rui Zhao. Shikra: Unleashing multimodal LLM’s referential dialogue magic. arXiv preprint arXiv:2306.15195, 2023.
- Kanzhi Cheng, Qiushi Sun, Yougang Chu, Fangzhi Xu, Yantao Li, Jianbing Zhang, and Zhiyong Wu. SeeClick: Harnessing GUI grounding for advanced visual GUI agents. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9313–9332, 2024. doi: 10.18653/v1/2024.acl-long.505.
- Filippos Christianos, Georgios Papoudakis, Thomas Coste, Jianye Hao, Jun Wang, and Kun Shao. Lightweight neural app control. In *International Conference on Learning Representations (ICLR)*, 2025.
- Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Samuel Stevens, Boshi Wang, Huan Sun, and Yu Su. Mind2Web: Towards a generalist agent for the web. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Number 57 in Monographs on Statistics and Applied Probability. Chapman & Hall, New York, 1993.

Boyu Gou, Ruohan Wang, Boyuan Zheng, Yanan Xie, Cheng Chang, Yiheng Shu, Huan Sun, and Yu Su. Navigating the digital world as humans do: Universal visual grounding for GUI agents. In *International Conference on Learning Representations (ICLR)*, 2025.

Tianrui Guan, Fuxiao Liu, Xiyang Wu, Ruiqi Xian, Zongxia Li, Xiaoyu Liu, Xijun Wang, Lichang Chen, Furong Huang, Yaser Yacoob, Dinesh Manocha, and Tianyi Zhou. HallusionBench: An advanced diagnostic suite for entangled language hallucination and visual illusion in large vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 14375–14385, 2024. doi: 10.1109/CVPR52733.2024.01363.

Wenyi Hong, Weihan Wang, Qingsong Lv, Jiazheng Xu, Wenmeng Yu, Junhui Ji, Yan Wang, Zihan Wang, Yuxiao Dong, Ming Ding, and Jie Tang. CogAgent: A visual language model for GUI agents. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 14281–14290, 2024. doi: 10.1109/CVPR52733.2024.01354.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022.

Hugging Face Transformers contributors. Qwen2-VL used to work with `inputs_embeds` instead of `input_ids`, but no more. GitHub issue #35463 (opened December 31, 2024) and pull request #35466, <https://github.com/huggingface/transformers/issues/35463>, 2024.

Kaixin Li, Ziyang Meng, Hongzhan Lin, Ziyang Luo, Yuchen Tian, Jing Ma, Zhiyong Huang, and Tat-Seng Chua. ScreenSpot-Pro: GUI grounding for professional high-resolution computer use. In *Proceedings of the 33rd ACM International Conference on Multimedia (MM)*, pages 8778–8786, 2025. doi: 10.1145/3746027.3755688.

Wei Li, William Bishop, Alice Li, Chris Rawles, Folawiyo Campbell-Ajala, Divya Tyamagundlu, and Oriana Riva. On the effects of data scale on UI control agents. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024.

Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating object hallucination in large vision-language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 292–305, 2023. doi: 10.18653/v1/2023.emnlp-main.20.

Kevin Qinghong Lin, Linjie Li, Difei Gao, Zhengyuan Yang, Shiwei Wu, Zechen Bai, Stan Weixian Lei, Lijuan Wang, and Mike Zheng Shou. ShowUI: One vision-language-action model for GUI visual agent. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 19498–19508, 2025. doi: 10.1109/CVPR52734.2025.01816.

Yadong Lu, Jianwei Yang, Yelong Shen, and Ahmed Awadallah. OmniParser for pure vision based GUI agent. arXiv preprint arXiv:2408.00203, 2024.

Zhengxi Lu, Yuxiang Chai, Yaxuan Guo, Xi Yin, Liang Liu, Hao Wang, Han Xiao, Shuai Ren, Pengxiang Zhao, Guangyi Liu, Guanqing Xiong, and Hongsheng Li. UI-R1: Enhancing efficient action prediction of GUI agents by reinforcement learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pages 17608–17616, 2026. doi: 10.1609/aaai.v40i21.38816.

Tiange Luo, Lajanugen Logeswaran, Justin Johnson, and Honglak Lee. Visual test-time scaling for GUI agent grounding. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025.

Xinbei Ma, Zhuosheng Zhang, and Hai Zhao. CoCo-Agent: A comprehensive cognitive MLLM agent for smartphone GUI automation. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.

416 Zhiliang Peng, Wenhui Wang, Li Dong, Yaru Hao, Shaohan Huang, Shuming Ma, and Furu Wei.
417 Kosmos-2: Grounding multimodal large language models to the world. In *International Confer-*
418 *ence on Learning Representations (ICLR)*, 2024.

419 Zehan Qi, Xiao Liu, Iat Long Iong, Hanyu Lai, Xueqiao Sun, Wenyi Zhao, Yu Yang, Xinyue Yang,
420 Jiadai Sun, Shuntian Yao, Tianjie Zhang, Wei Xu, Jie Tang, and Yuxiao Dong. WebRL: Training
421 LLM web agents via self-evolving online curriculum reinforcement learning. In *International*
422 *Conference on Learning Representations (ICLR)*, 2025.

423 Yujia Qin, Yining Ye, Junjie Fang, Haoming Wang, Shihao Liang, Shizuo Tian, Junda Zhang, Jia-
424 hao Li, Yunxin Li, Shijue Huang, Wanjuan Zhong, Kuanye Li, Jiale Yang, Yu Miao, Woyu Lin,
425 Longxiang Liu, Xu Jiang, Qianli Ma, Jingyu Li, Xiaojun Xiao, Kai Cai, Chuang Li, Yaowei
426 Zheng, Chaolin Jin, Chen Li, Xiao Zhou, Minchao Wang, Haoli Chen, Zhaojian Li, Haihua Yang,
427 Haifeng Liu, Feng Lin, Tao Peng, Xin Liu, and Guang Shi. UI-TARS: Pioneering automated GUI
428 interaction with native agents. arXiv preprint arXiv:2501.12326, 2025.

429 Christopher Rawles, Alice Li, Daniel Rodriguez, Oriana Riva, and Timothy Lillicrap. Android in
430 the wild: A large-scale dataset for Android device control. In *Advances in Neural Information*
431 *Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

432 Christopher Rawles, Sarah Clinckemaillie, Yifan Chang, Jonathan Waltz, Gabrielle Lau, Marybeth
433 Fair, Alice Li, William Bishop, Wei Li, Folawiyi Campbell-Ajala, Daniel Toyama, Robert Berry,
434 Divya Tyamagundlu, Timothy Lillicrap, and Oriana Riva. AndroidWorld: A dynamic benchmark-
435 ing environment for autonomous agents. In *International Conference on Learning Representa-*
436 *tions (ICLR)*, 2025.

437 Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: En-
438 hanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024. doi:
439 10.1016/j.neucom.2023.127063.

440 Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu,
441 Jialin Wang, Wenbin Ge, Yang Fan, Kai Dang, Mengfei Du, Xuancheng Ren, Rui Men, Dayi-
442 heng Liu, Chang Zhou, Jingren Zhou, and Junyang Lin. Qwen2-VL: Enhancing vision-language
443 model’s perception of the world at any resolution. arXiv preprint arXiv:2409.12191, 2024.

444 Yiqin Wang, Haoji Zhang, Jingqi Tian, and Yansong Tang. Ponder & Press: Advancing visual
445 GUI agent towards general computer control. In *Findings of the Association for Computational*
446 *Linguistics: ACL 2025*, pages 1461–1473, 2025. doi: 10.18653/v1/2025.findings-acl.76.

447 Qianhui Wu, Kanzhi Cheng, Rui Yang, Chaoyun Zhang, Jianwei Yang, Huiqiang Jiang, Jian Mu,
448 Baolin Peng, Bo Qiao, Reuben Tan, Si Qin, Lars Liden, Qingwei Lin, Huan Zhang, Tong Zhang,
449 Jianbing Zhang, Dongmei Zhang, and Jianfeng Gao. GUI-Actor: Coordinate-free visual ground-
450 ing for GUI agents. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025a.

451 Qingyuan Wu, Jianheng Liu, Jianye Hao, Jun Wang, and Kun Shao. Advancing autonomous
452 VLM agents via variational subgoal-conditioned reinforcement learning. arXiv preprint
453 arXiv:2502.07949, 2025b. Presented at the NeurIPS 2025 Workshop on Vision Language Models:
454 Challenges of Real World Deployment.

455 Zhiyong Wu, Zhenyu Wu, Fangzhi Xu, Yian Wang, Qiushi Sun, Chengyou Jia, Kanzhi Cheng,
456 Zichen Ding, Liheng Chen, Paul Pu Liang, and Yu Qiao. OS-ATLAS: A foundation action model
457 for generalist GUI agents. In *International Conference on Learning Representations (ICLR)*,
458 2025c.

459 Yiheng Xu, Zekun Wang, Junli Wang, Dunjie Lu, Tianbao Xie, Amrita Saha, Doyen Sahoo, Tao
460 Yu, and Caiming Xiong. Aguis: Unified pure vision agents for autonomous GUI interaction. In
461 *International Conference on Machine Learning (ICML)*, 2025.

462 Jianwei Yang, Hao Zhang, Feng Li, Xueyan Zou, Chunyuan Li, and Jianfeng Gao. Set-
463 of-mark prompting unleashes extraordinary visual grounding in GPT-4V. arXiv preprint
464 arXiv:2310.11441, 2023.

465 Rui Yang, Qianhui Wu, Zhaoyang Wang, Hanyang Chen, Ke Yang, Hao Cheng, Huaxiu Yao, Baolin
466 Peng, Huan Zhang, Jianfeng Gao, and Tong Zhang. GUI-Libra: Training native GUI agents
467 to reason and act with action-aware supervision and partially verifiable RL. arXiv preprint
468 arXiv:2602.22190, 2026.

469 Yuhao Yang, Yue Wang, Dongxu Li, Ziyang Luo, Bei Chen, Chao Huang, and Junnan Li. Aria-
470 UI: Visual grounding for GUI instructions. In *Findings of the Association for Computational
471 Linguistics: ACL 2025*, pages 22418–22433, 2025. doi: 10.18653/v1/2025.findings-acl.1152.

472 Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao.
473 ReAct: Synergizing reasoning and acting in language models. In *International Conference on
474 Learning Representations (ICLR)*, 2023.

475 Haoxuan You, Haotian Zhang, Zhe Gan, Xianzhi Du, Bowen Zhang, Zirui Wang, Liangliang Cao,
476 Shih-Fu Chang, and Yinfei Yang. Ferret: Refer and ground anything anywhere at any granularity.
477 In *International Conference on Learning Representations (ICLR)*, 2024a.

478 Keen You, Haotian Zhang, Eldon Schoop, Floris Weers, Amanda Swearngin, Jeffrey Nichols,
479 Yinfei Yang, and Zhe Gan. Ferret-UI: Grounded mobile UI understanding with multimodal
480 LLMs. In *European Conference on Computer Vision (ECCV)*, pages 240–255, 2024b. doi:
481 10.1007/978-3-031-73039-9_14.

482 Chi Zhang, Zhao Yang, Jiaxuan Liu, Yanda Li, Yucheng Han, Xin Chen, Zebiao Huang, Bin Fu, and
483 Gang Yu. AppAgent: Multimodal agents as smartphone users. In *Proceedings of the 2025 CHI
484 Conference on Human Factors in Computing Systems (CHI)*, pages 1–20, 2025. doi: 10.1145/
485 3706598.3713600.

486 Zhuosheng Zhang and Aston Zhang. You only look at screens: Multimodal chain-of-action agents.
487 In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 3132–3149, 2024.
488 doi: 10.18653/v1/2024.findings-acl.186.

489 Boyuan Zheng, Boyu Gou, Jihyung Kil, Huan Sun, and Yu Su. GPT-4V(ision) is a generalist web
490 agent, if grounded. In *International Conference on Machine Learning (ICML)*, 2024.

491 Shijie Zhou, Viet Dac Lai, Hao Tan, Jihyung Kil, Wanrong Zhu, Changyou Chen, and Ruiyi Zhang.
492 GUI-AIMA: Aligning intrinsic multimodal attention with a context anchor for GUI grounding.
493 arXiv preprint arXiv:2511.00810, 2025.

A Training and evaluation details

Data. AITW screenshots in the mirror we use are stored as raw RGB byte buffers without an image header; we decode them by matching the buffer length against seven observed resolutions (for example 540×1140 and 412×732) with a portrait-aspect fallback. Mind2Web screenshots are downsampled so that no image exceeds one megapixel. Examples are drawn from the AITW training stream in stored order, filtered to the mix’s action labels; the first n eligible steps form the training slice and the next 250 (200 for the control mix) form validation. Consecutive steps belong to the same episodes, so a validation slice spans a few dozen episodes and may share its boundary episode with training. Table 8 lists the class mix of each validation slice.

Prompts. Stage 2 prompts contain the screenshot, the line `Goal: {goal}`, and the instruction `Predict the action coordinate.` (C uses `Predict the action type and coordinate.` and is trained to answer with the action word before the coordinate; D-text prefixes `Action: {word}.`). D-token places `<|action_slot|>` at the start of the text; exactly one slot per sequence is asserted at every step.

Optimization. AdamW with learning rate 2×10^{-5} and weight decay 0.01, batch size 1, gradient clipping at 1.0, two epochs, a fixed per-epoch shuffle seeded by the run seed. The action-embedding table (D-hook, D-token) and the auxiliary head (B) share the base learning rate. Evaluation is greedy with at most 16 new tokens; the coordinate is parsed with a regular expression and an unparseable output counts as a miss at distance $\sqrt{2}$. Greedy evaluation is deterministic for a trained model; models in different tables that share a seed are separately trained copies (Table 9 and the intervention runs retrain the variant), and their numbers differ by training nondeterminism.

Compute. Everything runs on single NVIDIA L4 GPUs. A Stage-2 run at 1,200 training examples takes about 35 minutes; the full study, including retracted and re-run cells, cost about 130 GPU-hours.

Statistics. For each contrast we align the per-example distances of the two variants on every shared seed. The cluster bootstrap resamples the validation examples with replacement and keeps all seeds of a sampled example together, recomputes the mean difference 10,000 times, and reports the percentile interval and the two-sided bootstrap p . The seed-level test is a paired t -test on the per-seed means. The earlier (seed, example)-unit bootstrap and its permutation counterpart are reported in the released logs; their intervals are narrower by about a third and every star in this paper was recomputed on the cluster basis.

B Stage 1, Mind2Web, and per-class tables

Table 5: Stage 1 action-type macro-F1 (three seeds). On Mind2Web the instruction names the action, so the screenshot adds little; on AITW the instruction describes a goal and the screenshot is what separates the classes.

Features	Mind2Web (2 classes)	AITW (5 classes)
text only	0.597 ± 0.011	0.141 ± 0.032
text + screenshot	0.605 ± 0.016	0.555 ± 0.036
zeroed (sanity)	0.469 ± 0.000	0.110 ± 0.051
TF-IDF + logistic regression	0.622	0.168
majority class	0.472	0.140

Table 6: Multimodal-Mind2Web grounding control (1,000 training and 250 evaluation steps, three seeds, seed-level only). The target is the center of the gold element; hit@bbox is the fraction of predicted points inside the element. Both models are at floor.

Variant	hit@0.10	hit@0.25	hit@bbox	mean L2 ↓
A: flat	0.365 ± 0.031	0.645 ± 0.035	0.095 ± 0.018	0.234 ± 0.023
D-hook: additive	0.355 ± 0.032	0.664 ± 0.052	0.104 ± 0.034	0.233 ± 0.027

Table 7: Removing *type* events from the training slice only, scored on the identical headline validation slice (three seeds). Deltas are paired cluster bootstraps over examples; seed-level p in parentheses.

Variant	Training stream	hit@0.10	click hit@0.10	scroll hit@0.10	Δ hit@0.10 vs. A (same stream)
A: flat	with type	0.255 ± 0.026	0.294 ± 0.042	0.304 ± 0.022	
A: flat	type removed	0.297 ± 0.024	0.365 ± 0.039	0.275 ± 0.013	
B: aux. loss	with type	0.305 ± 0.015	0.373 ± 0.020	0.290 ± 0.013	
B: aux. loss	type removed	0.319 ± 0.034	0.404 ± 0.063	0.246 ± 0.066	+0.021 [-0.004, +0.045] (0.391)
D-hook: additive	with type	0.300 ± 0.035	0.351 ± 0.048	0.341 ± 0.045	
D-hook: additive	type removed	0.330 ± 0.037	0.393 ± 0.054	0.348 ± 0.000	+0.046 [+0.008, +0.084]* (-)

526 C Type events removed from training

527 D Full scaling table

528 E The earlier three-condition D-token test

529 Table 9 is the test that preceded Table 4: separately trained D-token models (five seeds), decoded
530 with the gold embedding, the cyclic wrong map, and the table zeroed. Its numbers differ from the
531 D-token rows of Table 4 by training nondeterminism and agree in every conclusion.

532 F The inputs_embeds failure in detail

533 Qwen2-VL locates the image tokens in `input_ids` to assign each visual token a temporal,
534 height, and width position for its three-dimensional rotary embedding. When a caller supplies
535 `inputs_embeds` and no `input_ids`, the model cannot perform that lookup; in the transformers
536 releases our runs used (pinned at 4.45 or later; the current 5.9 release behaves the same) the posi-
537 tions become a one-dimensional sequence index replicated across the three axes, so every image
538 token is positioned as text and the two-dimensional layout the grounding task depends on is lost. A
539 caller on that path must also merge the visual features into the embeddings itself, a second oppor-
540 tunity for silent error. We inferred the cause from the code and from the frozen-embedding control
541 described in Section 6.2 rather than by dumping position ids at the time. The same dependence on
542 `input_ids` was reported as a generation crash in December 2024 [Hugging Face Transformers con-
543 tributors, 2024]; the silent degradation when images are present is the case a fine-tuning practitioner
544 meets. In the set of runs where the failure appeared, the fixed additive variant scored 0.392 ± 0.043
545 hit@0.10 on `taps_and_swipes` against the flat baseline’s 0.390 ± 0.010 ; the re-run of that control
546 with per-example logging (Table 3) gives 0.363 ± 0.021 against 0.382 ± 0.015 .

547 G Where the model looks

548 We measured where the models look, with the caveats that the measure is head-averaged attention at
549 one layer, that attention is not information flow, and that teacher-forcing the gold answer places the
550 gold x coordinate in the prefix of the token that predicts y . Under gold conditioning both embedding
551 models put 10 to 11% of that token’s image attention within 0.10 of the target, against a chance
552 share of 2.6%; the last prompt token, which our earlier visualizations used, puts about 4% there
553 under every conditioning, so most of the localization we can see happens after the x coordinate is
554 in the prefix. [Flat-A row and free-running probe (model’s own answer teacher-forced) from the
555 third attention run.] The interventions move attention the same way they move accuracy on this seed

Table 8: Conditioning advantage vs. training-set size (AITW all_with_coords, 3 seeds per cell). Intervals are 95% cluster bootstraps over validation examples (all seeds of a sampled example kept together); stars from the same bootstrap. Absolute values are not comparable across rows because the validation slice moves with n ; its click/scroll/type counts are given per row. The study is descriptive: 24 contrasts, no multiplicity correction.

n_{train}	val mix (c/s/t)	A hit@0.10	B – A	D-hook – A	B – A (hit@0.25)	D-hook – A (hit@0.25)
300	170/33/47	0.139 ± 0.032	+0.067 [+0.031, +0.103]***	+0.069 [+0.041, +0.099]***	-0.055 [-0.092, -0.019]**	+0.067 [+0.033, +0.101]***
500	198/12/40	0.315 ± 0.059	+0.011 [-0.019, +0.040]	+0.001 [-0.032, +0.033]	+0.004 [-0.027, +0.033]	+0.059 [+0.025, +0.092]***
800	141/71/38	0.261 ± 0.020	-0.040 [-0.073, -0.008]*	+0.048 [+0.009, +0.087]*	-0.039 [-0.069, -0.008]*	+0.064 [+0.025, +0.103]***
1200	169/46/35	0.255 ± 0.026	+0.051 [+0.021, +0.081]***	+0.045 [+0.015, +0.079]**	+0.071 [+0.036, +0.107]***	+0.055 [+0.017, +0.092]**
2500	166/53/31	0.244 ± 0.037	-0.012 [-0.039, +0.015]	+0.037 [+0.009, +0.067]**	-0.008 [-0.043, +0.028]	+0.104 [+0.067, +0.143]***
5000	187/29/34	0.363 ± 0.009	+0.005 [-0.025, +0.037]	+0.043 [+0.011, +0.076]**	-0.023 [-0.051, +0.005]	+0.032 [+0.001, +0.064]*

Table 9: D-token causal-use test: the same trained model evaluated with the gold action embedding, a wrong (cyclically permuted) one, and a zeroed one (5 seeds, paired bootstrap over pooled examples).

Conditioning	hit@0.05	hit@0.10	hit@0.25	mean L2 ↓
gold	0.140 ± 0.059	0.236 ± 0.058	0.502 ± 0.045	0.386 ± 0.020
wrong	0.014 ± 0.004	0.079 ± 0.013	0.200 ± 0.027	0.728 ± 0.081
zero	0.074 ± 0.055	0.167 ± 0.066	0.416 ± 0.103	0.489 ± 0.059
gold – wrong		+0.157 [+0.114, +0.201]***	+0.302 [+0.246, +0.356]***	-0.342 [-0.375, -0.309]***
gold – zero		+0.069 [+0.038, +0.100]***	+0.086 [+0.056, +0.117]***	-0.103 [-0.121, -0.086]***

(Figure 3, Table 10): a wrong type roughly halves the target share for D-hook, and zeroing costs D-token both attention and accuracy. On this one seed zeroing left D-hook’s attention and accuracy unchanged, which the five-seed study in Section 6.1 shows to be within the noise of a small real cost. Table 10 gives the target-attention share for the last prompt token (the one that predicts the opening parenthesis), the token that predicts the first x digit, and the token that predicts the first y digit.

Table 10: Share of image attention within 0.10 of the target (chance 0.026) at three probe positions, seed 42, 120 screens, 3/4-depth layer, gold answer teacher-forced.

Model	Position	gold	wrong (cyclic)	wrong (click↔scroll)	zero
D-hook	last prompt token	0.040	0.036	0.035	0.036
D-hook	pre-x token	0.039	0.036	0.037	0.039
D-hook	pre-y token	0.111	0.064	0.049	0.119
D-token	last prompt token	0.038	0.034	0.034	0.035
D-token	pre-x token	0.038	0.034	0.033	0.035
D-token	pre-y token	0.095	0.080	0.079	0.079

H Qualitative examples




Figure 3: Same trained model, four conditionings, 120 validation screens, one seed per variant. Left: share of image attention within 0.10 of the target from the token that predicts the y coordinate, with the gold answer teacher-forced (3/4-depth layer, head-averaged); the dotted line is the share of image tokens within that radius. Right: greedy hit@0.10 under the same conditioning. Error bars are 95% paired-bootstrap intervals of the difference from gold over the 120 screens.




Figure 4: Predictions of the flat baseline (red) and D-hook (blue) against the target (green) on held-out AITW screens at the headline setting, one seed. The panels are a fixed spread of outcomes rather than a selection of wins: two rescues, a case where conditioning hurts, a case where both are right, a scroll, and a degenerate *type* example on which both models miss at distance $\sqrt{2}$. The title reports the base rates on this seed at $r = 0.10$: of 169 clicks, 21 are rescued by conditioning, 13 are hurt, 45 are correct under both models, and 90 are missed by both. In the rescue panels the baseline's prediction is off-screen and drawn clamped at the corner.