
Decoupling What from Where: A Mechanism Study of Action-Type Supervision for GUI Grounding in a Small Vision-Language Model

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 A GUI agent decides which action to take and where on the screen to take it; native
2 agent models emit both in one token stream. We ask what supervising the
3 action type does for the grounding half of that decision, using Qwen2-VL-2B
4 with LoRA on Android in the Wild and Multimodal-Mind2Web. A flat baseline
5 is compared with five ways of supplying the type under matched data, compute,
6 adapters, and decoding: an auxiliary loss, a hard-routed action word, an additive
7 learned embedding, a prepended learned token, and the type written into the
8 prompt. With five seeds, a cluster bootstrap over validation examples, and seed-
9 level paired tests, two results hold. On a training stream that mixes clicks and
10 scrolls with type events whose recorded target is a fixed off-screen point, the aux-
11 iliary loss and the additive embedding raise click grounding by 6 to 10 points over
12 the baseline, and the gain is not explained by the baseline emitting that off-screen
13 point; on a stream of taps and swipes the same mechanisms change nothing. The
14 prepended token, the mechanism we originally hypothesized, is no better than the
15 baseline. Interventions on the trained models show that both learned embeddings
16 are consulted at inference, with a difference that matters for deployment: with its
17 embedding zeroed, the additive model returns to the baseline’s accuracy, whereas
18 the prepended-token model falls below it. We also correct our own description of a
19 silent failure: injecting the conditioning through `inputs_embeds` makes Qwen2-
20 VL fall back to one-dimensional positions for image tokens, which produced a
21 nine-point deficit that vanished once `input_ids` were preserved.

22 1 Introduction

23 A GUI agent maps a screenshot and an instruction to a low-level action: an action type such as *click*,
24 *scroll*, or *type*, and, for spatial actions, a target location. Current native agent models (UI-TARS [Qin
25 et al., 2025], OS-Atlas [Wu et al., 2025c], SeeClick [Cheng et al., 2024], ShowUI [Lin et al., 2025])
26 produce both in a single autoregressive stream. The two decisions are different kinds of prediction:
27 the action type is a small classification problem that depends on the goal and the screen state, and
28 the location is a localization problem in which the model has to find the evidence on the screen that
29 supports the action and commit to a point. Several recent agents separate the two, predicting the
30 type first and the target conditioned on it [Ma et al., 2024, Christianos et al., 2025]. What has not
31 been measured is what the type supervision itself does to the grounding model, and which of the
32 many ways of delivering it matters.

33 This paper is a matched-compute study of that question in a small model. We fine-tune Qwen2-
34 VL-2B [Wang et al., 2024] with LoRA [Hu et al., 2022] to emit a coordinate string, and compare
35 a flat baseline with five conditioning mechanisms that receive the action type in different forms: as

an auxiliary classification loss, as a hard-routed action word in the output, as an additive learned embedding, as a prepended learned token, and as a word in the prompt. Each variant sees the same data in the same order and uses the same optimizer, adapters, and decoding. We started from the hypothesis that the prepended token would be the right mechanism, because it gives the model a dedicated, attendable position for the action type.

The results support type supervision and refute our hypothesis about the mechanism, and they turn out to depend on a property of the data that we did not appreciate at first. Android in the Wild records type events with a touch point off the screen, which our coordinate serializer clamps to the origin, so a training stream that mixes clicks, scrolls, and type events contains a class whose target is a fixed point. On that stream, the auxiliary loss and the additive embedding raise hit@0.10 by +0.064 and +0.054 over the flat baseline (five seeds, cluster bootstrap over examples, seed-level $p = 0.004$ and 0.028), almost entirely on clicks. On a stream of taps and swipes the same mechanisms are within noise of the baseline. The obvious explanation, that the baseline learns to emit the off-screen point for clicks and the conditioned models do not, accounts for only a small part of the gap: the baseline emits it on 4% of clicks, and excluding those cases barely moves the numbers. **[Removing type events from the training stream while scoring the identical validation examples: result of the no-type experiment.]** The prepended token is within noise of the baseline, and the type written as a word into the prompt **[D-text result]**.

Interventions on the trained models explain the ranking. Decoding either embedding model with a wrong type collapses it, and decoding with a zeroed or class-averaged embedding costs accuracy in both, so both learned embeddings are consulted at inference; the difference is what is left when the signal is removed. The zeroed additive model sits at the flat baseline’s level, while the zeroed prepended-token model falls seven points below it. An earlier implementation injected the conditioning through `inputs_embeds`; Qwen2-VL then falls back to one-dimensional positions for its image tokens, and that version scored nine points below the baseline, a result we initially read as evidence against conditioning.

Our contributions are (1) a matched-compute comparison of five action-type conditioning mechanisms against a flat baseline, with five seeds, a cluster bootstrap over examples, seed-level tests, and a control stream where conditioning should not help; (2) a diagnosis of where the gain comes from, including a training stream with the degenerate class removed; (3) interventions and attention measurements that separate mechanisms that use the type at inference from those that only use it in training; and (4) a corrected account of a positional-encoding failure that inverted a conclusion in a small model.

2 Related Work

GUI grounding and agent models. Mind2Web [Deng et al., 2023] and Android in the Wild [Rawles et al., 2023] supply the web and mobile episodes we train on; AndroidControl [Li et al., 2024] and AndroidWorld [Rawles et al., 2025] extend mobile evaluation. SeeClick [Cheng et al., 2024] argued that grounding is the bottleneck for visual GUI agents and introduced ScreenSpot; ScreenSpot-Pro [Li et al., 2025] raises the resolution. UGround [Gou et al., 2025] and OS-Atlas [Wu et al., 2025c] scale grounding pretraining, and Aguvis [Xu et al., 2025], UI-TARS [Qin et al., 2025], CogAgent [Hong et al., 2024], ScreenAI [Baechler et al., 2024], and Ferret-UI [You et al., 2024b] train agents that emit action and location together, the design our flat baseline follows. GUI-Actor [Wu et al., 2025a] replaces coordinate text with an attention-based action head, UI-R1 [Lu et al., 2026] rewards the action type and the argument separately, RegionFocus [Luo et al., 2025] zooms at test time, and OmniParser [Lu et al., 2024] and Set-of-Mark prompting [Yang et al., 2023] externalize grounding into a parser.

Type-first agents. CoCo-Agent [Ma et al., 2024] decomposes AITW action prediction into the action type followed by the target conditioned on it, in one stream; that is our hard-routing variant. LiMAC [Christianos et al., 2025] places a small action-type transformer ahead of a LoRA-tuned Qwen2-VL-2B on AITW and AndroidControl and selects click targets conditioned on the predicted type, using screen elements rather than pixels. SeeAct [Zheng et al., 2024], Ponder & Press [Wang et al., 2025], and Aria-UI [Yang et al., 2025] separate text planning from pixel grounding. GUI-Libra [Yang et al., 2026] reweights reasoning against action tokens during fine-tuning and GUI-AIMA [Zhou et al., 2025] adds an anchor token for grounding. These systems establish that type-




Figure 1: Two-stage pipeline. Stage 1 classifies the action type from mean-pooled frozen features. Stage 2 grounds the action to a coordinate string, and the action type enters it through one of the mechanisms in the bottom row (D-text, the type as a word in the prompt, is not drawn). All mechanisms share data, compute, adapters, and decoding. D-slot is the retracted variant whose `inputs_embeds` injection bypassed the position computation.

first pipelines work; none of them holds data, compute, and decoding fixed across conditioning mechanisms, and none intervenes on the trained model to ask whether the type signal is used, which is what this paper does.

Conditioning, faithfulness, and the backbone. SayCan [Ahn et al., 2022], ReAct [Yao et al., 2023], and RT-2 [Brohan et al., 2023] are the reference points for factored, exposed, and token-level action decisions; AutoUI [Zhang and Zhang, 2024], AppAgent [Zhang et al., 2025], DigiRL [Bai et al., 2024], WebRL [Qi et al., 2025], and VSC-RL [Wu et al., 2025b] improve mobile agents by imitation or reinforcement learning without changing how the type enters the grounding computation. Our auxiliary-loss variant is a multi-task objective in the sense of Caruana [1997]. Kosmos-2 [Peng et al., 2024], Shikra [Chen et al., 2023], and Ferret [You et al., 2024a] train VLMs to refer to image regions, and POPE [Li et al., 2023] and HallusionBench [Guan et al., 2024] ask whether an output is supported by the image; our attention measurements ask that of a coordinate. Qwen2-VL [Wang et al., 2024] assigns three-dimensional rotary positions (extending RoPE [Su et al., 2024]) to visual tokens from `input_ids`; Section 6.2 describes what happens when conditioning bypasses that path.

3 Method

3.1 Action taxonomy and grounding target

Both datasets are mapped onto a canonical eight-class action space (*click*, *double-click*, *type*, *scroll*, *drag*, *hotkey*, *wait*, *finished*). Mind2Web’s CLICK, SELECT, and HOVER map to *click* and TYPE to *type*. AITW stores gestures as touch and lift points; we split them into taps (*click*) and swipes (*scroll*) by the normalized distance between the two, with threshold 0.04. The grounding target is the touch point normalized to $[0, 1]^2$, serialized as (x, y) with integers on a 0–999 grid so every target has the same token length. AITW records non-touch actions such as typing with the touch point $(-1, -1)$; the serializer clamps this to $(0, 0)$, so every *type* event in the training stream has the same target at the screen origin, and its distance to the raw target is exactly $\sqrt{2}$ whatever the model predicts. This detail matters for the analysis in Section 5.2.

3.2 Stage 1: which action

Stage 1 runs the frozen Qwen2-VL-2B-Instruct backbone over the screenshot and instruction, mean-pools the final hidden states over unmasked tokens into a 1536-dimensional vector, and trains a three-layer MLP ($1536 \rightarrow 1024 \rightarrow 256 \rightarrow 8$, dropout 0.1) with class-weighted cross-entropy, AdamW (learning rate 10^{-4} , weight decay 0.01), batch size 128, for eight epochs, keeping the checkpoint with the best validation macro-F1. We compare features from text only, text plus screenshot, and a sanity variant with zeroed features.

122 3.3 Stage 2: where

123 Stage 2 fine-tunes the same backbone with LoRA adapters (rank 16, $\alpha = 32$, dropout 0.05) on the
 124 query, key, value, and output projections, 4.4M trainable parameters out of 2.2B. The prompt is the
 125 screenshot, `Goal: {goal}`, and a fixed instruction to predict the action coordinate; the answer is
 126 the coordinate string. With s the screenshot, g the goal, a the action type, and $y_{1:T}$ the coordinate
 127 tokens, every variant minimizes

$$\mathcal{L}_{\text{coord}} = - \sum_{t=1}^T \log p_{\theta}(y_t | y_{<t}, s, g, c), \quad (1)$$

128 masked to the answer tokens, with c the variant-specific conditioning.

129 **A, flat.** $c = \emptyset$: the joint-decoding design of native agent models, reduced to the grounding sub-
 130 problem.

131 **B, auxiliary loss.** A linear head h_{ϕ} on the mean-pooled last hidden state $\bar{z}$ predicts the action type
 132 during training, $\mathcal{L}_B = \mathcal{L}_{\text{coord}} + \lambda \text{CE}(h_{\phi}(\bar{z}), a)$ with $\lambda = 1$. Nothing about the type is provided at
 133 inference, so any gain is a training effect.

134 **C, hard routing.** The model is trained to emit the action word before the coordinate, `click (x,`
 135 `y)`, as in CoCo-Agent’s conditional action prediction. At test time the word for the gold (or Stage-1)
 136 type is forced as a decoding prefix.

137 **D-hook, additive embedding.** A learned table $E \in \mathbb{R}^{8 \times d}$ holds one vector per class; a forward hook
 138 on the embedding layer adds $E[a]$ to every text and image token embedding. E is zero-initialized,
 139 so D-hook starts as an exact copy of A.

140 **D-token, prepended action token.** A new token `<|action_slot|>` is added to the vocabulary
 141 and placed at the start of the text prompt, so it occupies a real position in `input_ids`; a forward
 142 hook replaces its embedding with $E[a]$, initialized from $\mathcal{N}(0, 0.02^2)$. This is the architecture we
 143 originally hypothesized: a full-norm, attendable representation of the type.

144 **D-text, type as a word.** The prompt begins with `Action: click.` (or the gold type’s word),
 145 using the pretrained word embeddings and no new parameters. This is the obvious baseline for the
 146 two learned-embedding variants.

147 For C, D-hook, D-token, and D-text the gold type is used in training and at evaluation unless stated
 148 otherwise; Section 5.4 closes the loop with Stage-1 predictions. All variants share the data order, the
 149 optimizer (AdamW, learning rate 2×10^{-5} , weight decay 0.01, batch size 1, gradient clipping at 1.0,
 150 two epochs), and greedy decoding with 16 new tokens. No hyperparameter was tuned per variant;
 151 Section 7 returns to this.

152 4 Experimental setup

153 **Data.** Stage 1 uses 7,775 Multimodal-Mind2Web training steps evaluated on the 1,338-step
 154 `test_task` split, and 5,000 AITW steps from the General subset evaluated on 1,000 held-out
 155 steps. Stage 2 uses two AITW streams built by filtering the training split in its stored order.
 156 `all_with_coords` keeps taps, swipes, and type events; at 1,200 training examples its validation
 157 slice of 250 steps holds 169 clicks, 46 scrolls, and 35 type events. `taps_and_swipes` keeps only
 158 taps and swipes (1,000 training and 200 validation steps); both land anywhere on the screen. The
 159 AITW stream is deterministic, so for a given mix and training size every seed and every variant
 160 is evaluated on the same examples in the same order, which is what makes paired tests possible.
 161 The validation slice begins where the training slice ends; consecutive steps come from the same
 162 episodes, and absolute numbers are not comparable across training sizes. A second control uses
 163 Multimodal-Mind2Web (1,000 training and 250 evaluation steps), where the target is the center of
 164 the gold element.

165 **Metrics.** `hit@r` is the fraction of predictions within normalized Euclidean distance r of the target,
 166 for $r \in \{0.05, 0.10, 0.25\}$; `hit@0.10` is the headline. We also report the mean normalized distance.
 167 Unparseable outputs count as misses at distance $\sqrt{2}$; only hard routing produces them (2 to 8% of
 168 outputs), and its numbers are reported both ways. Per-class numbers are computed on the same basis
 169 as the overall ones.

Table 1: Stage 2 grounding on AITW all_with_coords ($n_{\text{train}}=1200$, $n_{\text{val}}=250$, 5 seeds). Mean $\pm$ std over seeds; Δ with 95% paired-bootstrap CI over pooled (seed, example) units; * $p < .05$, ** $p < .01$, *** $p < .001$.

Variant	hit@0.05	hit@0.10	hit@0.25	mean L2 $\downarrow$	Δ hit@0.10 vs. A [cluster CI]	seed-level p
A: flat	0.130 $\pm$ 0.039	0.229 $\pm$ 0.043	0.499 $\pm$ 0.029	0.406 $\pm$ 0.028		
B: aux. loss	0.178 $\pm$ 0.025	0.293 $\pm$ 0.021	0.582 $\pm$ 0.026	0.365 $\pm$ 0.004	+0.064 [+0.036, +0.093]***	0.004
C: hard routing	0.170 $\pm$ 0.022	0.264 $\pm$ 0.045	0.538 $\pm$ 0.059	0.421 $\pm$ 0.016	+0.035 [+0.001, +0.069]*	0.189
D-hook: additive	0.170 $\pm$ 0.042	0.282 $\pm$ 0.041	0.563 $\pm$ 0.016	0.370 $\pm$ 0.010	+0.054 [+0.022, +0.086]***	0.028
D-token: prepended	0.144 $\pm$ 0.056	0.238 $\pm$ 0.049	0.504 $\pm$ 0.055	0.386 $\pm$ 0.021	+0.009 [-0.015, +0.034]	0.465

Statistics. The headline setting uses five seeds (42 to 46); every other cell uses three (42 to 44). We report the mean and sample standard deviation over seeds. Every run logs a per-example distance on a shared evaluation set, and every seed of every variant shares the per-seed data order, so comparisons are paired at two levels. Our primary interval is a cluster bootstrap that resamples validation examples with replacement and keeps all seeds of a sampled example together (10,000 resamples, 95% percentile interval, two-sided bootstrap p [Efron and Tibshirani, 1993]); resampling (seed, example) units separately would treat the seeds of one example as independent and gives intervals about a third narrower. Because the seeds share a training trajectory only through the data order, we also report a paired t -test on the five (or three) per-seed means, which has little power but cannot be fooled by within-seed correlation. We treat a difference as established when both agree; stars in tables come from the cluster bootstrap, and the seed-level p is printed beside them. No correction is made for the number of contrasts; the scaling study in Section 5.5 is presented as descriptive for that reason.

5 Results

5.1 Stage 1: the screenshot decides the action type on AITW

Stage 1 (Table 5 in the appendix, three seeds) motivates the choice of dataset. On Mind2Web the instruction usually names the action, a bag-of-words classifier reaches 0.622 macro-F1, and the screenshot adds +0.009; on AITW the instruction is a goal (“open a new private window”), text-only features sit at the majority floor (0.141), and the screenshot adds +0.414. Predicting the type is a vision problem on AITW, so the Stage-2 study is run there.

5.2 Type supervision improves click grounding on the mixed stream

Table 1 is the main result: six variants on all_with_coords at 1,200 training examples, five seeds each. The auxiliary loss (B) and the additive embedding (D-hook) beat the flat baseline on every metric, by +0.064 and +0.054 hit@0.10 with cluster intervals that exclude zero and seed-level p of 0.004 and 0.028, and they are indistinguishable from each other (D-hook – B: -0.010 , cluster CI $[-0.043, +0.024]$, seed $p = 0.56$). Hard routing (C) gains +0.035 with a cluster interval that barely excludes zero and a seed-level p of 0.19, and loses on mean distance, partly because forcing the action word makes 2 to 8% of its outputs unparseable; scored on parsed outputs only, its click rate is unchanged (0.356 vs. 0.350). The prepended token (D-token), our hypothesized mechanism, is within noise of the baseline on every hit rate (+0.009, cluster CI $[-0.015, +0.034]$) and below B by -0.055 (cluster CI $[-0.088, -0.024]$, seed $p = 0.016$). [D-text row: the type as a word in the prompt scores X.] Seed variance is large: the baseline ranges from 0.169 to 0.280 across seeds, which is why we report five seeds and paired statistics rather than a best run. The ordering with the first three seeds was the same; the two added seeds lowered every variant and widened the gap between D-token and the two working mechanisms.

Table 2 shows where the gain lives and tests the simplest explanation for it. The improvement is a click improvement: B and C gain ten points on clicks, D-hook six, and C pays for its click gain with a sixteen-point loss on scrolls because forcing the word *scroll* into the output derails the gesture; D-hook is the only variant that improves both classes. Since the only difference between this stream and the control is the type class with its fixed target, one might expect the flat baseline’s click errors to be emissions of that target that the conditioned models avoid. They mostly are not. The baseline emits (0, 0) for 4.3% of clicks (13% on its worst seed), the auxiliary-loss model for 2.4%, and the three conditioned-at-inference models never; excluding those predictions moves the baseline’s click

Table 2: Where the gain lives. Per-class hit@0.10 on the headline mix (five seeds, same basis as Table 1), the fraction of click predictions that are the clamped type target (0, 0), and click hit@0.10 after excluding those predictions. Every model scores zero on *type* because the raw target is off-screen.

Variant	click	scroll	type	clicks sent to (0, 0)	click, excluding those
A: flat	0.256 $\pm$ 0.062	0.304 $\pm$ 0.034	0.000	4.3% (13.0% on the worst seed)	0.266
B: aux. loss	0.357 $\pm$ 0.026	0.278 $\pm$ 0.018	0.000	2.4%	0.366
C: hard routing	0.350 $\pm$ 0.061	0.148 $\pm$ 0.073	0.000	0.0%	0.350
D-hook: additive	0.321 $\pm$ 0.064	0.357 $\pm$ 0.039	0.000	0.0%	0.321
D-token: prepended	0.269 $\pm$ 0.069	0.304 $\pm$ 0.090	0.000	0.0%	0.269

Table 3: Control: AITW taps_and_swipes ($n_{\text{train}}=1000$, $n_{\text{val}}=200$, 3 seeds), where action type is not spatially informative.

Variant	hit@0.05	hit@0.10	hit@0.25	mean L2 $\downarrow$	Δ hit@0.10 vs. A [cluster CI]	seed-level p
A: flat	0.243 $\pm$ 0.020	0.382 $\pm$ 0.015	0.720 $\pm$ 0.017	0.185 $\pm$ 0.003		
B: aux. loss	0.225 $\pm$ 0.015	0.368 $\pm$ 0.012	0.730 $\pm$ 0.039	0.173 $\pm$ 0.011	-0.013 [-0.038, +0.012]	0.456
C: hard routing	0.217 $\pm$ 0.016	0.325 $\pm$ 0.040	0.635 $\pm$ 0.064	0.271 $\pm$ 0.068	-0.057 [-0.108, -0.007]*	0.215
D-hook: additive	0.235 $\pm$ 0.022	0.363 $\pm$ 0.021	0.725 $\pm$ 0.043	0.182 $\pm$ 0.009	-0.018 [-0.052, +0.015]	0.334

rate from 0.256 to 0.266 while B stays at 0.366 and D-hook at 0.321. The baseline’s remaining click predictions are also shifted toward the left edge by 0.10 on average on the seed whose predictions we logged (D-hook’s by 0.02), so the presence of the degenerate class degrades the baseline’s click grounding more broadly than by exact emissions. [No-type experiment, Table 7: removing type events from the training slice while scoring the identical validation examples gives ...] [sentence stating whether the gain needs the class present in training].

5.3 Control: no gain on taps and swipes

On taps_and_swipes (Table 3), whose training stream has no degenerate class, B and D-hook are within noise of the baseline (-0.013 and -0.018 hit@0.10, cluster intervals $[-0.038, +0.012]$ and $[-0.052, +0.015]$), so the data admit harms of up to four points as well as no effect, and C is worse by -0.057 (cluster CI $[-0.108, -0.007]$, seed $p = 0.22$) with a large increase in mean distance. The Multimodal-Mind2Web control sits at floor (a tenth of predicted points fall inside the gold element for both A and D-hook, three seeds) and is uninformative in either direction; we report it in the appendix and do not lean on it.

5.4 End to end: predicted types instead of gold

The tables above condition on the gold action type. To close the loop we train D-hook, extract frozen features for its training and validation examples, fit the Stage-1 MLP on the three classes present for a fixed 200 full-batch epochs with no checkpoint selection, and evaluate D-hook twice on the same examples: with gold types and with Stage-1 predictions (three seeds). Stage 1 reaches 0.843 accuracy and 0.795 macro-F1. The oracle pipeline scores 0.292 hit@0.10, the predicted-type pipeline 0.271, and the flat baseline on the same seeds 0.255. Only the oracle-to-predicted gap is established ($+0.021$, cluster CI $[+0.003, +0.041]$, seed $p = 0.07$); the predicted pipeline’s margin over the baseline ($+0.016$, cluster CI $[-0.013, +0.047]$, seed $p = 0.44$) is not, and we report it as descriptive. Classifier error removes more than half of the conditioning gain in this setting.

5.5 Scaling: a descriptive study

Figure 2 extends the comparison to $n \in \{300, 500, 800, 1200, 2500, 5000\}$ with three seeds at every size. The cluster intervals for D-hook – A exclude zero at five of the six sizes on hit@0.10 (all six on hit@0.25), and its point estimate at 5,000 examples ($+0.043$) is the same size as at 1,200; B – A excludes zero at 300 and 1,200, is negative at 800, and is within noise at 500, 2,500, and 5,000. With three seeds the seed-level tests reach $p < 0.05$ only at 300 and 5,000 for D-hook and at 300 and 1,200 for B, and a direct D-hook – B contrast is within noise at every size except 800 and 2,500. We read the study as consistent with the additive embedding’s advantage persisting with data and the auxiliary loss’s advantage not, and no stronger than that; an earlier single-seed version of the




Figure 2: Conditioning advantage over the flat baseline as a function of training-set size on `all_with_coords`, three seeds per cell, 95% cluster-bootstrap intervals over examples. Each size is scored on a different validation slice with a different class mix (Table 8), so only within-size differences are meaningful, and 24 contrasts are shown without correction.

Table 4: Interventions on the trained conditioned models: the same model is decoded with the gold action embedding, a wrong one (cyclic permutation; `click→type`), a wrong one that swaps only clicks and scrolls, the table zeroed, and every row replaced by the class mean (uninformative, in-distribution norm). Five seeds $\times$ 250 examples; Δ = gold – condition with 95% cluster-bootstrap CI over examples and the seed-level paired p . Flat A on the same examples is the reference row.

Model	Conditioning	hit@0.10	click hit@0.10	scroll hit@0.10	Δ gold – cond. (hit@0.10)	seed p
D-hook: additive	gold type	0.275 ± 0.040	0.310 ± 0.060	0.357 ± 0.052		
	wrong (cyclic)	0.054 ± 0.009	0.006 ± 0.000	0.270 ± 0.050	$+0.222$ [$+0.174, +0.270$]***	0.000
	wrong (click↔scroll)	0.099 ± 0.056	0.073 ± 0.087	0.270 ± 0.050	$+0.176$ [$+0.133, +0.221$]***	0.003
	zeroed	0.248 ± 0.032	0.288 ± 0.053	0.291 ± 0.050	$+0.027$ [$-0.006, +0.060$]	0.224
	class mean	0.206 ± 0.058	0.217 ± 0.099	0.322 ± 0.064	$+0.070$ [$+0.034, +0.106$]***	0.133
	flat A, same examples	0.229 ± 0.043	0.256	0.304	$+0.046$ [$+0.017, +0.077$]**	0.032
D-token: prepended	gold type	0.276 ± 0.034	0.325 ± 0.026	0.304 ± 0.132		
	wrong (cyclic)	0.084 ± 0.014	0.037 ± 0.027	0.319 ± 0.066	$+0.192$ [$+0.139, +0.244$]***	0.004
	wrong (click↔scroll)	0.144 ± 0.036	0.126 ± 0.048	0.319 ± 0.066	$+0.132$ [$+0.077, +0.187$]***	0.069
	zeroed	0.180 ± 0.088	0.187 ± 0.134	0.290 ± 0.033	$+0.096$ [$+0.053, +0.140$]***	0.307
	class mean	0.175 ± 0.055	0.178 ± 0.085	0.297 ± 0.033	$+0.101$ [$+0.057, +0.147$]***	0.182
	flat A, same examples	0.229 ± 0.043	0.256	0.304	$+0.021$ [$-0.007, +0.051$]	0.047

curve had suggested that all conditioning gains shrink with data, which three seeds do not support. Figure 4 in the appendix shows what the numbers look like on individual screens.

6 Mechanism

6.1 Interventions on both embedding mechanisms

D-token’s failure to help could mean that the model ignores the slot, and D-hook’s success could be a training effect that leaves a vector the model never reads. To test both on equal terms we retrain each at the headline setting (five seeds) and decode every trained model under five conditionings on the full validation slice: the gold type; a wrong type from a cyclic permutation of the classes present (click to type, type to scroll, scroll to click); a second wrong map that swaps only clicks and scrolls and so never routes a groundable example to the degenerate class; the table zeroed; and every row replaced by the mean of the rows of the classes present, which is uninformative but has an in-distribution norm. Table 4 reports the results with the flat baseline on the same examples as the reference. (An earlier three-condition version of this test on separately trained D-token models, Table 9 in the appendix, agrees.) The learned rows of D-hook are small: each has norm 0.04 against a mean token-embedding norm of 0.58, whereas D-token’s rows have norm 0.78. Small as they are, the model reads them. A wrong type collapses D-hook from 0.275 to 0.054 under the cyclic map and to 0.099 under the click-scroll map (cluster intervals [$+0.174, +0.270$] and [$+0.133, +0.221$], seed $p < 0.001$ and 0.003), so the collapse is not an artifact of routing clicks to the degenerate class. The class-mean vector costs seven points ($+0.070$, cluster CI [$+0.034, +0.106$], seed $p = 0.13$). Zeroing costs about three, and that difference is not established ($+0.027$, cluster CI [$-0.006, +0.060$], seed

266 $p = 0.22$): the zeroed model sits at the level of the flat baseline (+0.019 over A, not established).
 267 The single-seed attention run in Section ??, on which an earlier version of this paper based the claim
 268 that the embedding could be removed without loss, is consistent with these intervals and was simply
 269 underpowered. The additive embedding therefore behaves as a switch at inference: with the right
 270 type the model gains six points over the baseline, without a signal it returns to the baseline’s level,
 271 and with a wrong or blended signal it collapses.

272 D-token behaves differently in one respect that matters for deployment. Every perturbation collapses
 273 it, including the in-distribution class-mean condition (+0.101, cluster CI [+0.057, +0.147]), so its
 274 sensitivity is not an artifact of feeding it an out-of-distribution zero vector; and the zeroed model
 275 is worse than the flat baseline by -0.075 (cluster CI $[-0.112, -0.039]$), whereas the zeroed D-
 276 hook is not. Adaptation with a prepended token makes the model depend on the token; adaptation
 277 with an additive vector does not, even though the vector is used. Per-class rows show that in both
 278 mechanisms the conditioning acts almost entirely on clicks: scroll accuracy is nearly unchanged
 279 under every condition. [re-check D-token numbers after seeds 45 and 46 land] Attention measured
 280 at the token that predicts the y coordinate moves the same way as accuracy under these interventions
 281 on a single seed (Appendix G); we treat those numbers as a consistency check, not as evidence about
 282 faithfulness on their own.

283 6.2 A positional-encoding failure that inverted a conclusion

284 Our first implementation of the prepended token replaced the slot’s embedding by building
 285 inputs_embeds and passing that to the model instead of input_ids. On taps_and_swipes that
 286 variant scored 0.298 ± 0.026 hit@0.10 against 0.390 ± 0.010 for the flat baseline in the same set of
 287 runs, and a diagnostic variant with the same slot but a frozen random embedding scored the same
 288 (0.290), which pointed at the injection path rather than the learned vector. Qwen2-VL derives its
 289 three-dimensional rotary positions from input_ids; given inputs_embeds alone, the transformers
 290 implementation falls back to sequential one-dimensional positions replicated across the three axes,
 291 so the image tokens lose their two-dimensional layout entirely (details, version, and a related library
 292 report [Hugging Face Transformers contributors, 2024] in Appendix F). Keeping input_ids intact
 293 and modifying embeddings through a forward hook removed the deficit. We record this because the
 294 failure is silent, because we initially read it as evidence against conditioning, and because an earlier
 295 version of this paper described its mechanism incorrectly as a shift of the visual positions.

296 7 Limitations

297 The study is small: one 2B backbone, LoRA adapters, validation slices of 200 to 250 steps, one learn-
 298 ing rate and one adapter configuration for every variant, and one L4 per run. The absolute grounding
 299 numbers are far below the field’s leaderboards, per-seed variance is large, and the conclusions rest
 300 on paired statistics that we report at two levels because neither is clean: the cluster bootstrap ignores
 301 that seeds share a data order, and the seed-level test has five or three units. The refutation of the
 302 prepended token is conditional on that untuned configuration; a stronger D-token could exist under
 303 a schedule that trains the fresh token faster, and we did not train a control that keeps the slot but
 304 freezes its embedding. The headline stream’s only spatially informative class is one whose target
 305 is a fixed off-screen point, so the claim is about adaptation on a mixed action stream that contains
 306 such a class, not about a general spatial prior; the type-removed experiment bounds what the class
 307 contributes. The end-to-end result is descriptive. The scaling study scores each size on a different
 308 validation slice, and its 24 contrasts are uncorrected. The attention analysis uses one seed, one layer,
 309 and a probe token chosen after observing that the last prompt token does not localize. Mind2Web is
 310 at floor here and tells us nothing. We do not evaluate on ScreenSpot or on multi-step task success,
 311 where the compounding-error motivation of the introduction would have to be tested directly. Code,
 312 split indices, and prompts will be released.

313 8 Conclusion

314 Supervising the action type during adaptation of a small grounding model helps when the training
 315 stream mixes groundable actions with a class whose target is degenerate, and the help is concentrated
 316 on clicks; on a stream of taps and swipes it does nothing. The mechanism ranking is stable across

317 five seeds and two levels of paired inference: an auxiliary loss and an additive embedding capture
 318 the gain, the additive embedding keeps it across training sizes from 300 to 5,000 examples, and
 319 the prepended learned token we set out to validate is consulted by the model and improves nothing
 320 over a baseline that never saw a type. Both learned embeddings are read at inference and both
 321 collapse under a wrong type; what distinguishes them is that a model adapted with the additive
 322 vector still grounds at the baseline’s level when the vector is removed, while a model adapted with
 323 the prepended token does not. For practitioners adapting small VLMs as GUI agents on mixed action
 324 streams, the usable lessons are to supervise the action type, to prefer conditioning that degrades to the
 325 baseline rather than below it when the type is missing or wrong, and to keep `input_ids` on the path
 326 whenever a conditioning scheme touches the embeddings of a model with multimodal positional
 327 encoding.

328 References

- 329 Michael Ahn, Anthony Brohan, Noah Brown, Yevgen Chebotar, Omar Cortes, Byron David, Chelsea
 330 Finn, Chuyuan Fu, Keerthana Gopalakrishnan, Karol Hausman, Alex Herzog, Daniel Ho, Jasmine
 331 Hsu, Julian Ibarz, Brian Ichter, Alex Irpan, Eric Jang, Rosario Jauregui Ruano, Kyle Jeffrey, Sally
 332 Jesmonth, Nikhil J Joshi, Ryan Julian, Dmitry Kalashnikov, Yuheng Kuang, Kuang-Huei Lee,
 333 Sergey Levine, Yao Lu, Linda Luu, Carolina Parada, Peter Pastor, Jornell Quiambao, Kanishka
 334 Rao, Jarek Rettinghouse, Diego Reyes, Pierre Sermanet, Nicolas Sievers, Clayton Tan, Alexander
 335 Toshev, Vincent Vanhoucke, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Mengyuan Yan, and Andy
 336 Zeng. Do as I can, not as I say: Grounding language in robotic affordances. In *Conference on*
 337 *Robot Learning (CoRL)*, 2022.
- 338 Gilles Baechler, Srinivas Sunkara, Maria Wang, Fedir Zubach, Hassan Mansoor, Vincent Etter, Vic-
 339 tor Cărbune, Jason Lin, Jindong Chen, and Abhanshu Sharma. ScreenAI: A vision-language
 340 model for UI and infographics understanding. In *Proceedings of the Thirty-Third International*
 341 *Joint Conference on Artificial Intelligence (IJCAI)*, 2024.
- 342 Hao Bai, Yifei Zhou, Mert Cemri, Jiayi Pan, Alane Suhr, Sergey Levine, and Aviral Kumar. DigiRL:
 343 Training in-the-wild device-control agents with autonomous reinforcement learning. In *Advances*
 344 *in Neural Information Processing Systems (NeurIPS)*, 2024.
- 345 Anthony Brohan, Noah Brown, Justice Carbajal, Yevgen Chebotar, Xi Chen, Krzysztof Choro-
 346 manski, Tianli Ding, Danny Driess, Avinava Dubey, Chelsea Finn, Pete Florence, Chuyuan Fu,
 347 Montse Gonzalez Arenas, Keerthana Gopalakrishnan, Kehang Han, Karol Hausman, Alexander
 348 Herzog, Jasmine Hsu, Brian Ichter, Alex Irpan, Nikhil Joshi, Ryan Julian, Dmitry Kalashnikov,
 349 Yuheng Kuang, Isabel Leal, Lisa Lee, Tsang-Wei Edward Lee, Sergey Levine, Yao Lu, Hen-
 350 ryk Michalewski, Igor Mordatch, Karl Pertsch, Kanishka Rao, Krista Reymann, Michael Ryoo,
 351 Grecia Salazar, Pannag Sanketi, Pierre Sermanet, Jaspiar Singh, Anikait Singh, Radu Soricut,
 352 Huang Tran, Vincent Vanhoucke, Quan Vuong, Ayzaan Wahid, Stefan Welker, Paul Wohlhart,
 353 Jialin Wu, Fei Xia, Ted Xiao, Peng Xu, Sichun Xu, Tianhe Yu, and Brianna Zitkovich. RT-
 354 2: Vision-language-action models transfer web knowledge to robotic control. In *Conference on*
 355 *Robot Learning (CoRL)*, 2023.
- 356 Rich Caruana. Multitask learning. *Machine Learning*, 28(1):41–75, 1997. doi: 10.1023/A:
 357 1007379606734.
- 358 Keqin Chen, Zhao Zhang, Weili Zeng, Richong Zhang, Feng Zhu, and Rui Zhao. Shikra: Unleashing
 359 multimodal LLM’s referential dialogue magic. arXiv preprint arXiv:2306.15195, 2023.
- 360 Kanzhi Cheng, Qiushi Sun, Yougang Chu, Fangzhi Xu, Yantao Li, Jianbing Zhang, and Zhiyong
 361 Wu. SeeClick: Harnessing GUI grounding for advanced visual GUI agents. In *Proceedings of the*
 362 *62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*,
 363 pages 9313–9332, 2024. doi: 10.18653/v1/2024.acl-long.505.
- 364 Filippos Christianos, Georgios Papoudakis, Thomas Coste, Jianye Hao, Jun Wang, and Kun Shao.
 365 Lightweight neural app control. In *International Conference on Learning Representations (ICLR)*,
 366 2025.

367 Xiang Deng, Yu Gu, Boyuan Zheng, Shijie Chen, Samuel Stevens, Boshi Wang, Huan Sun, and
368 Yu Su. Mind2Web: Towards a generalist agent for the web. In *Advances in Neural Information*
369 *Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

370 Bradley Efron and Robert J. Tibshirani. *An Introduction to the Bootstrap*. Number 57 in Monographs
371 on Statistics and Applied Probability. Chapman & Hall, New York, 1993.

372 Boyu Gou, Ruohan Wang, Boyuan Zheng, Yanan Xie, Cheng Chang, Yiheng Shu, Huan Sun, and
373 Yu Su. Navigating the digital world as humans do: Universal visual grounding for GUI agents. In
374 *International Conference on Learning Representations (ICLR)*, 2025.

375 Tianrui Guan, Fuxiao Liu, Xiyang Wu, Ruiqi Xian, Zongxia Li, Xiaoyu Liu, Xijun Wang, Lichang
376 Chen, Furong Huang, Yaser Yacoob, Dinesh Manocha, and Tianyi Zhou. HallusionBench: An
377 advanced diagnostic suite for entangled language hallucination and visual illusion in large vision-
378 language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern*
379 *Recognition (CVPR)*, pages 14375–14385, 2024. doi: 10.1109/CVPR52733.2024.01363.

380 Wenyi Hong, Weihang Wang, Qingsong Lv, Jiazhen Xu, Wenmeng Yu, Junhui Ji, Yan Wang, Zihan
381 Wang, Yuxiao Dong, Ming Ding, and Jie Tang. CogAgent: A visual language model for GUI
382 agents. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*
383 *(CVPR)*, pages 14281–14290, 2024. doi: 10.1109/CVPR52733.2024.01354.

384 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang,
385 and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Con-*
386 *ference on Learning Representations (ICLR)*, 2022.

387 Hugging Face Transformers contributors. Qwen2-VL used to work with inputs_embeds instead
388 of input_ids, but no more. GitHub issue #35463 (opened December 31, 2024) and pull request
389 #35466, <https://github.com/huggingface/transformers/issues/35463>, 2024.

390 Kaixin Li, Ziyang Meng, Hongzhan Lin, Ziyang Luo, Yuchen Tian, Jing Ma, Zhiyong Huang, and
391 Tat-Seng Chua. ScreenSpot-Pro: GUI grounding for professional high-resolution computer use.
392 In *Proceedings of the 33rd ACM International Conference on Multimedia (MM)*, pages 8778–
393 8786, 2025. doi: 10.1145/3746027.3755688.

394 Wei Li, William Bishop, Alice Li, Chris Rawles, Folawiyo Campbell-Ajala, Divya Tyamagundlu,
395 and Oriana Riva. On the effects of data scale on UI control agents. In *Advances in Neural*
396 *Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2024.

397 Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating
398 object hallucination in large vision-language models. In *Proceedings of the 2023 Conference*
399 *on Empirical Methods in Natural Language Processing (EMNLP)*, pages 292–305, 2023. doi:
400 10.18653/v1/2023.emnlp-main.20.

401 Kevin Qinghong Lin, Linjie Li, Difei Gao, Zhengyuan Yang, Shiwei Wu, Zechen Bai, Stan Weix-
402 ian Lei, Lijuan Wang, and Mike Zheng Shou. ShowUI: One vision-language-action model for
403 GUI visual agent. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern*
404 *Recognition (CVPR)*, pages 19498–19508, 2025. doi: 10.1109/CVPR52734.2025.01816.

405 Yadong Lu, Jianwei Yang, Yelong Shen, and Ahmed Awadallah. OmniParser for pure vision based
406 GUI agent. arXiv preprint arXiv:2408.00203, 2024.

407 Zhengxi Lu, Yuxiang Chai, Yaxuan Guo, Xi Yin, Liang Liu, Hao Wang, Han Xiao, Shuai
408 Ren, Pengxiang Zhao, Guangyi Liu, Guanjing Xiong, and Hongsheng Li. UI-R1: Enhanc-
409 ing efficient action prediction of GUI agents by reinforcement learning. In *Proceedings of*
410 *the AAAI Conference on Artificial Intelligence*, volume 40, pages 17608–17616, 2026. doi:
411 10.1609/aaai.v40i21.38816.

412 Tiange Luo, Lajanugen Logeswaran, Justin Johnson, and Honglak Lee. Visual test-time scaling for
413 GUI agent grounding. In *Proceedings of the IEEE/CVF International Conference on Computer*
414 *Vision (ICCV)*, 2025.

415 Xinbei Ma, Zhuosheng Zhang, and Hai Zhao. CoCo-Agent: A comprehensive cognitive MLLM
416 agent for smartphone GUI automation. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024.
417

418 Zhiliang Peng, Wenhui Wang, Li Dong, Yaru Hao, Shaohan Huang, Shuming Ma, and Furu Wei.
419 Kosmos-2: Grounding multimodal large language models to the world. In *International Conference on Learning Representations (ICLR)*, 2024.
420

421 Zehan Qi, Xiao Liu, Iat Long Iong, Hanyu Lai, Xueqiao Sun, Wenyi Zhao, Yu Yang, Xinyue Yang,
422 Jiadai Sun, Shuntian Yao, Tianjie Zhang, Wei Xu, Jie Tang, and Yuxiao Dong. WebRL: Training
423 LLM web agents via self-evolving online curriculum reinforcement learning. In *International
424 Conference on Learning Representations (ICLR)*, 2025.

425 Yujia Qin, Yining Ye, Junjie Fang, Haoming Wang, Shihao Liang, Shizuo Tian, Junda Zhang, Jia-
426 hao Li, Yunxin Li, Shijue Huang, Wanjuan Zhong, Kuanye Li, Jiale Yang, Yu Miao, Woyu Lin,
427 Longxiang Liu, Xu Jiang, Qianli Ma, Jingyu Li, Xiaojun Xiao, Kai Cai, Chuang Li, Yaowei
428 Zheng, Chaolin Jin, Chen Li, Xiao Zhou, Minchao Wang, Haoli Chen, Zhaojian Li, Haihua Yang,
429 Haifeng Liu, Feng Lin, Tao Peng, Xin Liu, and Guang Shi. UI-TARS: Pioneering automated GUI
430 interaction with native agents. arXiv preprint arXiv:2501.12326, 2025.

431 Christopher Rawles, Alice Li, Daniel Rodriguez, Oriana Riva, and Timothy Lillicrap. Android in
432 the wild: A large-scale dataset for Android device control. In *Advances in Neural Information
433 Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.

434 Christopher Rawles, Sarah Clinckemaulle, Yifan Chang, Jonathan Waltz, Gabrielle Lau, Marybeth
435 Fair, Alice Li, William Bishop, Wei Li, Folawiyi Campbell-Ajala, Daniel Toyama, Robert Berry,
436 Divya Tyamagundlu, Timothy Lillicrap, and Oriana Riva. AndroidWorld: A dynamic benchmark-
437 ing environment for autonomous agents. In *International Conference on Learning Representa-
438 tions (ICLR)*, 2025.

439 Jianlin Su, Yu Lu, Shengfeng Pan, Ahmed Murtadha, Bo Wen, and Yunfeng Liu. RoFormer: En-
440 hanced transformer with rotary position embedding. *Neurocomputing*, 568:127063, 2024. doi:
441 10.1016/j.neucom.2023.127063.

442 Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu,
443 Jialin Wang, Wenbin Ge, Yang Fan, Kai Dang, Mengfei Du, Xuancheng Ren, Rui Men, Dayi-
444 heng Liu, Chang Zhou, Jingren Zhou, and Junyang Lin. Qwen2-VL: Enhancing vision-language
445 model’s perception of the world at any resolution. arXiv preprint arXiv:2409.12191, 2024.

446 Yiqin Wang, Haoji Zhang, Jingqi Tian, and Yansong Tang. Ponder & Press: Advancing visual
447 GUI agent towards general computer control. In *Findings of the Association for Computational
448 Linguistics: ACL 2025*, pages 1461–1473, 2025. doi: 10.18653/v1/2025.findings-acl.76.

449 Qianhui Wu, Kanzhi Cheng, Rui Yang, Chaoyun Zhang, Jianwei Yang, Huiqiang Jiang, Jian Mu,
450 Baolin Peng, Bo Qiao, Reuben Tan, Si Qin, Lars Liden, Qingwei Lin, Huan Zhang, Tong Zhang,
451 Jianbing Zhang, Dongmei Zhang, and Jianfeng Gao. GUI-Actor: Coordinate-free visual ground-
452 ing for GUI agents. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025a.

453 Qingyuan Wu, Jianheng Liu, Jianye Hao, Jun Wang, and Kun Shao. Advancing autonomous
454 VLM agents via variational subgoal-conditioned reinforcement learning. arXiv preprint
455 arXiv:2502.07949, 2025b. Presented at the NeurIPS 2025 Workshop on Vision Language Models:
456 Challenges of Real World Deployment.

457 Zhiyong Wu, Zhenyu Wu, Fangzhi Xu, Yian Wang, Qiushi Sun, Chengyou Jia, Kanzhi Cheng,
458 Zichen Ding, Liheng Chen, Paul Pu Liang, and Yu Qiao. OS-ATLAS: A foundation action model
459 for generalist GUI agents. In *International Conference on Learning Representations (ICLR)*,
460 2025c.

461 Yiheng Xu, Zekun Wang, Junli Wang, Dunjie Lu, Tianbao Xie, Amrita Saha, Doyen Sahoo, Tao
462 Yu, and Caiming Xiong. Aguvis: Unified pure vision agents for autonomous GUI interaction. In
463 *International Conference on Machine Learning (ICML)*, 2025.

464 Jianwei Yang, Hao Zhang, Feng Li, Xueyan Zou, Chunyuan Li, and Jianfeng Gao. Set-
 465 of-mark prompting unleashes extraordinary visual grounding in GPT-4V. arXiv preprint
 466 arXiv:2310.11441, 2023.

467 Rui Yang, Qianhui Wu, Zhaoyang Wang, Hanyang Chen, Ke Yang, Hao Cheng, Huaxiu Yao, Baolin
 468 Peng, Huan Zhang, Jianfeng Gao, and Tong Zhang. GUI-Libra: Training native GUI agents
 469 to reason and act with action-aware supervision and partially verifiable RL. arXiv preprint
 470 arXiv:2602.22190, 2026.

471 Yuhao Yang, Yue Wang, Dongxu Li, Ziyang Luo, Bei Chen, Chao Huang, and Junnan Li. Aria-
 472 UI: Visual grounding for GUI instructions. In *Findings of the Association for Computational*
 473 *Linguistics: ACL 2025*, pages 22418–22433, 2025. doi: 10.18653/v1/2025.findings-acl.1152.

474 Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao.
 475 ReAct: Synergizing reasoning and acting in language models. In *International Conference on*
 476 *Learning Representations (ICLR)*, 2023.

477 Haoxuan You, Haotian Zhang, Zhe Gan, Xianzhi Du, Bowen Zhang, Zirui Wang, Liangliang Cao,
 478 Shih-Fu Chang, and Yinfei Yang. Ferret: Refer and ground anything anywhere at any granularity.
 479 In *International Conference on Learning Representations (ICLR)*, 2024a.

480 Keen You, Haotian Zhang, Eldon Schoop, Floris Weers, Amanda Swearngin, Jeffrey Nichols,
 481 Yinfei Yang, and Zhe Gan. Ferret-UI: Grounded mobile UI understanding with multimodal
 482 LLMs. In *European Conference on Computer Vision (ECCV)*, pages 240–255, 2024b. doi:
 483 10.1007/978-3-031-73039-9_14.

484 Chi Zhang, Zhao Yang, Jiaxuan Liu, Yanda Li, Yucheng Han, Xin Chen, Zebiao Huang, Bin Fu, and
 485 Gang Yu. AppAgent: Multimodal agents as smartphone users. In *Proceedings of the 2025 CHI*
 486 *Conference on Human Factors in Computing Systems (CHI)*, pages 1–20, 2025. doi: 10.1145/
 487 3706598.3713600.

488 Zhuosheng Zhang and Aston Zhang. You only look at screens: Multimodal chain-of-action agents.
 489 In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 3132–3149, 2024.
 490 doi: 10.18653/v1/2024.findings-acl.186.

491 Boyuan Zheng, Boyu Gou, Jihyung Kil, Huan Sun, and Yu Su. GPT-4V(ision) is a generalist web
 492 agent, if grounded. In *International Conference on Machine Learning (ICML)*, 2024.

493 Shijie Zhou, Viet Dac Lai, Hao Tan, Jihyung Kil, Wanrong Zhu, Changyou Chen, and Ruiyi Zhang.
 494 GUI-AIMA: Aligning intrinsic multimodal attention with a context anchor for GUI grounding.
 495 arXiv preprint arXiv:2511.00810, 2025.

496 A Training and evaluation details

497 **Data.** AITW screenshots in the mirror we use are stored as raw RGB byte buffers without an
 498 image header; we decode them by matching the buffer length against seven observed resolutions
 499 (for example 540×1140 and 412×732) with a portrait-aspect fallback. Mind2Web screenshots are
 500 downsampled so that no image exceeds one megapixel. Examples are drawn from the AITW training
 501 stream in stored order, filtered to the mix’s action labels; the first n eligible steps form the training
 502 slice and the next 250 (200 for the control mix) form validation. Consecutive steps belong to the
 503 same episodes, so a validation slice spans a few dozen episodes and may share its boundary episode
 504 with training. Table 8 lists the class mix of each validation slice.

505 **Prompts.** Stage 2 prompts contain the screenshot, the line `Goal: {goal}`, and the instruction
 506 `Predict the action coordinate.` (C uses `Predict the action type and coordinate.`
 507 and is trained to answer with the action word before the coordinate; D-text prefixes `Action:`
 508 `{word}.`). D-token places `<|action_slot|>` at the start of the text; exactly one slot per sequence
 509 is asserted at every step.

510 **Optimization.** AdamW with learning rate 2×10^{-5} and weight decay 0.01, batch size 1, gradient
 511 clipping at 1.0, two epochs, a fixed per-epoch shuffle seeded by the run seed. The action-embedding
 512 table (D-hook, D-token) and the auxiliary head (B) share the base learning rate. Evaluation is greedy
 513 with at most 16 new tokens; the coordinate is parsed with a regular expression and an unparseable
 514 output counts as a miss at distance $\sqrt{2}$. Greedy evaluation is deterministic for a trained model;
 515 models in different tables that share a seed are separately trained copies (Table 9 and the intervention
 516 runs retrain the variant), and their numbers differ by training nondeterminism.

517 **Compute.** Everything runs on single NVIDIA L4 GPUs. A Stage-2 run at 1,200 training examples
 518 takes about 35 minutes; the full study, including retracted and re-run cells, cost about 130 GPU-
 519 hours.

520 **Statistics.** For each contrast we align the per-example distances of the two variants on every shared
 521 seed. The cluster bootstrap resamples the validation examples with replacement and keeps all seeds
 522 of a sampled example together, recomputes the mean difference 10,000 times, and reports the per-
 523 centile interval and the two-sided bootstrap p . The seed-level test is a paired t -test on the per-seed
 524 means. The earlier (seed, example)-unit bootstrap and its permutation counterpart are reported in
 525 the released logs; their intervals are narrower by about a third and every star in this paper was
 526 recomputed on the cluster basis.

527 B Stage 1, Mind2Web, and per-class tables

Table 5: Stage 1 action-type macro-F1 (three seeds). On Mind2Web the instruction names the action, so the screenshot adds little; on AITW the instruction describes a goal and the screenshot is what separates the classes.

Features	Mind2Web (2 classes)	AITW (5 classes)
text only	0.597 ± 0.011	0.141 ± 0.032
text + screenshot	0.605 ± 0.016	0.555 ± 0.036
zeroed (sanity)	0.469 ± 0.000	0.110 ± 0.051
TF-IDF + logistic regression	0.622	0.168
majority class	0.472	0.140

Table 6: Multimodal-Mind2Web grounding control (1,000 training and 250 evaluation steps, three seeds, seed-level only). The target is the center of the gold element; hit@bbox is the fraction of predicted points inside the element. Both models are at floor.

Variant	hit@0.10	hit@0.25	hit@bbox	mean L2 ↓
A: flat	0.365 ± 0.031	0.645 ± 0.035	0.095 ± 0.018	0.234 ± 0.023
D-hook: additive	0.355 ± 0.032	0.664 ± 0.052	0.104 ± 0.034	0.233 ± 0.027

Table 7: Removing *type* events from the training slice only, scored on the identical headline validation slice (three seeds). Deltas are paired cluster bootstraps over examples; seed-level *p* in parentheses.

Variant	Training stream	hit@0.10	click hit@0.10	scroll hit@0.10	Δ hit@0.10 vs. A (same stream)
A: flat	with type	0.255 ± 0.026	0.294 ± 0.042	0.304 ± 0.022	
A: flat	type removed	0.292	0.355	0.283	
B: aux. loss	with type	0.305 ± 0.015	0.373 ± 0.020	0.290 ± 0.013	
D-hook: additive	with type	0.300 ± 0.035	0.351 ± 0.048	0.341 ± 0.045	

528 C Type events removed from training

529 D Full scaling table

530 E The earlier three-condition D-token test

531 Table 9 is the test that preceded Table 4: separately trained D-token models (five seeds), decoded
532 with the gold embedding, the cyclic wrong map, and the table zeroed. Its numbers differ from the
533 D-token rows of Table 4 by training nondeterminism and agree in every conclusion.

534 F The inputs_embeds failure in detail

535 Qwen2-VL locates the image tokens in `input_ids` to assign each visual token a temporal,
536 height, and width position for its three-dimensional rotary embedding. When a caller supplies
537 `inputs_embeds` and no `input_ids`, the model cannot perform that lookup; in the transformers
538 releases our runs used (pinned at 4.45 or later; the current 5.9 release behaves the same) the posi-
539 tions become a one-dimensional sequence index replicated across the three axes, so every image
540 token is positioned as text and the two-dimensional layout the grounding task depends on is lost. A
541 caller on that path must also merge the visual features into the embeddings itself, a second oppor-
542 tunity for silent error. We inferred the cause from the code and from the frozen-embedding control
543 described in Section 6.2 rather than by dumping position ids at the time. The same dependence on
544 `input_ids` was reported as a generation crash in December 2024 [Hugging Face Transformers con-
545 tributors, 2024]; the silent degradation when images are present is the case a fine-tuning practitioner
546 meets. In the set of runs where the failure appeared, the fixed additive variant scored 0.392 ± 0.043
547 hit@0.10 on `taps_and_swipes` against the flat baseline’s 0.390 ± 0.010 ; the re-run of that control
548 with per-example logging (Table 3) gives 0.363 ± 0.021 against 0.382 ± 0.015 .

549 G Where the model looks

550 We measured where the models look, with the caveats that the measure is head-averaged attention at
551 one layer, that attention is not information flow, and that teacher-forcing the gold answer places the
552 gold x coordinate in the prefix of the token that predicts y. Under gold conditioning both embedding
553 models put 10 to 11% of that token’s image attention within 0.10 of the target, against a chance
554 share of 2.6%; the last prompt token, which our earlier visualizations used, puts about 4% there
555 under every conditioning, so most of the localization we can see happens after the x coordinate is
556 in the prefix. [Flat-A row and free-running probe (model’s own answer teacher-forced) from the
557 third attention run.] The interventions move attention the same way they move accuracy on this seed
558 (Figure 3, Table 10): a wrong type roughly halves the target share for D-hook, and zeroing costs
559 D-token both attention and accuracy. On this one seed zeroing left D-hook’s attention and accuracy

Table 8: Conditioning advantage vs. training-set size (AITW all_with_coords, 3 seeds per cell). Intervals are 95% cluster bootstraps over validation examples (all seeds of a sampled example kept together); stars from the same bootstrap. Absolute values are not comparable across rows because the validation slice moves with n ; its click/scroll/type counts are given per row. The study is descriptive: 24 contrasts, no multiplicity correction.

n_{train}	val mix (c/s/t)	A hit@0.10	B - A	D-hook - A	B - A (hit@0.25)	D-hook - A (hit@0.25)
300	170/33/47	0.139 $\pm$ 0.032	+0.067 [+0.031, +0.103]***	+0.069 [+0.041, +0.099]***	-0.055 [-0.092, -0.019]**	+0.067 [+0.033, +0.101]***
500	198/12/40	0.315 $\pm$ 0.059	+0.011 [-0.019, +0.040]	+0.001 [-0.032, +0.033]	+0.004 [-0.027, +0.033]	+0.059 [+0.025, +0.092]***
800	141/71/38	0.261 $\pm$ 0.020	-0.040 [-0.073, -0.008]*	+0.048 [+0.009, +0.087]*	-0.039 [-0.069, -0.008]*	+0.064 [+0.025, +0.103]***
1200	169/46/35	0.255 $\pm$ 0.026	+0.051 [+0.021, +0.081]***	+0.045 [+0.015, +0.079]**	+0.071 [+0.036, +0.107]***	+0.055 [+0.017, +0.092]**
2500	166/53/31	0.244 $\pm$ 0.037	-0.012 [-0.039, +0.015]	+0.037 [+0.009, +0.067]**	-0.008 [-0.043, +0.028]	+0.104 [+0.067, +0.143]***
5000	187/29/34	0.363 $\pm$ 0.009	+0.005 [-0.025, +0.037]	+0.043 [+0.011, +0.076]**	-0.023 [-0.051, +0.005]	+0.032 [+0.001, +0.064]*

Table 9: D-token causal-use test: the same trained model evaluated with the gold action embedding, a wrong (cyclically permuted) one, and a zeroed one (5 seeds, paired bootstrap over pooled examples).

Conditioning	hit@0.05	hit@0.10	hit@0.25	mean L2 $\downarrow$
gold	0.140 $\pm$ 0.059	0.236 $\pm$ 0.058	0.502 $\pm$ 0.045	0.386 $\pm$ 0.020
wrong	0.014 $\pm$ 0.004	0.079 $\pm$ 0.013	0.200 $\pm$ 0.027	0.728 $\pm$ 0.081
zero	0.074 $\pm$ 0.055	0.167 $\pm$ 0.066	0.416 $\pm$ 0.103	0.489 $\pm$ 0.059
gold - wrong		+0.157 [+0.114, +0.201]***	+0.302 [+0.246, +0.356]***	-0.342 [-0.375, -0.309]***
gold - zero		+0.069 [+0.038, +0.100]***	+0.086 [+0.056, +0.117]***	-0.103 [-0.121, -0.086]***

560 unchanged, which the five-seed study in Section 6.1 shows to be within the noise of a small real cost.
561 Table 10 gives the target-attention share for the last prompt token (the one that predicts the opening
562 parenthesis), the token that predicts the first x digit, and the token that predicts the first y digit.

Table 10: Share of image attention within 0.10 of the target (chance 0.026) at three probe positions, seed 42, 120 screens, 3/4-depth layer, gold answer teacher-forced.

Model	Position	gold	wrong (cyclic)	wrong (click $\leftrightarrow$ scroll)	zero
D-hook	last prompt token	0.040	0.036	0.035	0.036
D-hook	pre-x token	0.039	0.036	0.037	0.039
D-hook	pre-y token	0.111	0.064	0.049	0.119
D-token	last prompt token	0.038	0.034	0.034	0.035
D-token	pre-x token	0.038	0.034	0.033	0.035
D-token	pre-y token	0.095	0.080	0.079	0.079

563 H Qualitative examples




Figure 3: Same trained model, four conditionings, 120 validation screens, one seed per variant. Left: share of image attention within 0.10 of the target from the token that predicts the y coordinate, with the gold answer teacher-forced (3/4-depth layer, head-averaged); the dotted line is the share of image tokens within that radius. Right: greedy hit@0.10 under the same conditioning. Error bars are 95% paired-bootstrap intervals of the difference from gold over the 120 screens.




Figure 4: Predictions of the flat baseline (red) and D-hook (blue) against the target (green) on held-out AITW screens at the headline setting, one seed. The panels are a fixed spread of outcomes rather than a selection of wins: two rescues, a case where conditioning hurts, a case where both are right, a scroll, and a degenerate *type* example on which both models miss at distance $\sqrt{2}$. The title reports the base rates on this seed at $r = 0.10$: of 169 clicks, 21 are rescued by conditioning, 13 are hurt, 45 are correct under both models, and 90 are missed by both. In the rescue panels the baseline's prediction is off-screen and drawn clamped at the corner.