
When Does Action-Type Supervision Help GUI Grounding?

Separating Real Effects from Evaluation Artifacts

Anonymous Author(s)

Affiliation

Address

email

Abstract

GUI agents must predict an action type (*click*, *scroll*, *type*, ...) and a spatial target from a screenshot and an instruction. We study whether explicitly factoring these decisions helps a 2B-parameter agent, via a matched-compute comparison of five conditioning mechanisms on Qwen2-VL: a flat baseline (A), an auxiliary action-type loss (B), hard routing (C), an additive embedding hook (D-hook), and the prepended learned action embedding our original proposal advocated (D-token). Our first evaluation of this question was wrong in three independent ways, each of which we diagnose, quantify, and repair: (i) injecting conditioning via `inputs_embeds` silently bypasses Qwen2-VL’s M-RoPE position path and manufactured an 8σ *false negative*; (ii) train/validation splits taken as consecutive slices of one episode stream leaked shared task goals (24% of validation steps at the headline setting, up to 85% at scale) while looking clean at the episode level (2%); and (iii) a paired bootstrap pooling (seed, example) predictions as independent units was anti-conservative, rejecting a true clustered null 21.2% of the time at nominal 5% (an episode-clustered replacement achieves 5.1%). We then rebuilt the data pipeline around frozen goal-disjoint splits with a test region no phase of the project had ever touched, prespecified four contrasts, and evaluated once. **Every mechanistic conclusion of the original study reversed**, and the statistical repair alone would not have caught most of them. The auxiliary loss that the repaired statistics still endorsed on historical data ($B-A +0.051 [+0.020, +0.081]$, $p = 0.0043$, versus the invalid pooled $p = 0.0002$) is null on the untouched test set ($+0.011 [-0.005, +0.029]$, Holm $p = 0.309$); the prepended action embedding our proposal advocated and the course-era version reported as *refuted* in fact beats it ($D\text{-token}-B +0.035 [+0.017, +0.054]$, Holm $p = 0.0027$) — though D-token reads the ground-truth action type at test time and B reads nothing, so that contrast is oracle-versus-deployable; the largest *relative* gain moves to *scroll* ($2.5\times$), while click still supplies 61% of the absolute gain; the control condition designed to show no effect shows one, and is therefore invalid; and the deployable predicted-type pipeline is null ($+0.015 [-0.004, +0.034]$, Holm $p = 0.309$) because Stage 1 accuracy falls from 0.795 on the frozen validation split to 0.708 on the untouched test region, absorbing 76% of the oracle gain and leaving no detectable benefit. What survives is narrow and honest: *oracle* action-type conditioning improves grounding ($D\text{-hook}-A +0.057 [+0.035, +0.079]$, Holm $p = \leq 0.0004$), and nothing we can actually deploy does. We release the frozen splits, per-example predictions, and the clustered-inference code. A study can be multi-seed, ablated,

controlled, and apparently significant at the pooled $p=0.0002$ we now retract, and still have every one of its mechanistic claims move under a clean split and a calibrated test.

1 Introduction

A GUI agent maps a screenshot and a natural-language instruction to a low-level action: a type (*click*, *scroll*, *type*, *drag*, ...) and, for spatial actions, a pixel-level target. State-of-the-art native agents (UI-TARS [13], OS-Atlas [11], SeeClick [3], ShowUI [6]) emit both in one autoregressive stream, entangling the two decisions. A documented failure mode follows: commit to the wrong action type, then ground *coherently* within the wrong semantic frame – scrolling past a visible target instead of clicking it. This motivated our original hypothesis, formulated as an architecture bet: a learned action-type embedding, prepended to the instruction, should let a grounding model condition on *what* before deciding *where*, and should help most in low data.

This paper reports what happened when we tested that bet carefully, and it has two intertwined subjects. The first is the scientific question: *when does action-type supervision change GUI grounding, and through which mechanism?* The second is methodological: *which implementation and evaluation choices can create a false appearance of benefit or harm in studies of exactly this kind?* We came to the second subject involuntarily – three separate artifacts in our own pipeline each produced a wrong conclusion at some point during the project – and we believe the post-mortem is at least as useful to practitioners as the object-level result.

Concretely, we compare five matched-compute Stage-2 conditioning mechanisms on Qwen2-VL-2B [10] with LoRA, over AITW [9] with a Mind2Web [4] cross-dataset control, preceded by a Stage-1 action-type classifier study. Our contributions:

- **A confirmatory answer, and it is mostly negative.** On an untouched, goal-disjoint test set: auxiliary action-type supervision does not measurably help; *oracle* inference-time conditioning does (+0.057 [+0.035, +0.079]); the deployable predicted-type pipeline does not, because the action-type classifier drops from 0.795 to 0.708 accuracy between the two held-out regions; and the proportionally largest effect moves to *scroll* rather than the click disambiguation our exploratory analysis reported, though click’s higher frequency still makes it 61% of the absolute gain.
- **Three quantified evaluation pitfalls.** (1) A silent M-RoPE / inputs_embeds interaction that fabricated an 8σ false negative; (2) episode-stream train/validation leakage that is invisible at the episode level (2% of validation steps) but large at the shared-task-goal level (24%–85%); (3) pseudo-replicated significance testing whose anti-conservativeness we measure on synthetic clustered nulls (21.2% type-I error at nominal 5%).
- **A repaired, releasable protocol.** Pinned-revision, goal-disjoint frozen train/validation/test manifests with an untouched test region beyond every historical data access; episode-clustered paired inference with per-seed effects and validated calibration; and a prespecified confirmatory analysis committed before test access. All artifacts, per-example predictions, and analysis code are released.
- **A measured cost of getting evaluation wrong.** Because we ran the same ablation twice — once under the flawed protocol, once under the repaired one — we can report what the flaws actually bought: not inflated p -values alone, but a reversal of which mechanism wins, which action class carries the gain, whether our own hypothesis was refuted, whether the pipeline is deployable, and whether the control holds (Section 7.4).

2 Related Work

End-to-end GUI agents. The dominant paradigm treats action type, arguments, and coordinates as one autoregressive output. UI-TARS [13] exemplifies this with a unified action space across plat-




Figure 1: Stage 1 action-type classifier macro-F1 on Mind2Web (2 active classes) and AITW (5 active classes). Vision is essential on AITW (+0.414 over text-only) but redundant on Mind2Web (+0.009), because AITW instructions describe *goals* rather than physically specifying the action.

forms and explicit thought traces. OS-Atlas [11] scales the same paradigm with 13M cross-platform GUI grounding instances and shows that action-label harmonization across web/Android/desktop is nontrivial. ShowUI [6] builds a compact VLA around Qwen2-VL with interleaved vision-language-action streaming. Our flat baseline (Variant A) replicates the joint-decoding paradigm of these systems.

Modular grounding. SeeClick [3] argues that grounding is the central bottleneck in visual GUI agents. UGround / SeeAct-V [5] separates planning from grounding at the text-vs-pixel boundary. GUI-Actor [14] replaces coordinate-token regression with an attention-based action head driven by a dedicated <ACTOR> token plus a verifier. Our work studies a different module boundary: between *action-type prediction* and *spatial grounding*, not between planning text and coordinates.

Reward decomposition. UI-R1 [7] applies RL fine-tuning with a rule-based reward that separates action-type accuracy from action-argument accuracy on AndroidControl. It is the closest training analog: type and location as separable optimization targets; we use supervised conditioning rather than RL. RegionFocus [8] refines grounding at inference; we change the training-time conditioning pathway. SayCan [1] and ReAct [12] provide the modular-decision precedent, and RT-2 [2] the unified token-level action generation paradigm we critically examine.

Evaluation integrity. Our methodological findings echo, in a GUI setting, long-standing concerns about test-set contamination and dependence-blind inference: grouped data require grouped splits and cluster-aware uncertainty, and repeated measurements of the same example must not be pooled as independent evidence. Mind2Web’s own metric design (element accuracy vs. operation F1 vs. step success) [4] reflects the type/location decomposition we study.

3 Data and Splits

Sources. AITW [9] (HuggingFace mirror `cjfcst/AITW_General`, revision pinned to `5c0dc713`) is the primary Stage-2 benchmark; Mind2Web [4] provides the Stage-1 comparison and a cross-dataset control. Two AITW data mixes matter: `all_with_coords` (tap, swipe, type) and `taps_and_swipes`, intended as a control in which action type carries little information about *where* — an assumption Section 7.2 shows to be false. Targets are normalized to $[0, 1]^2$ and emitted as "`x`,

110 y) " on a $[0, 999]$ grid, with loss masked to coordinate tokens. Mirror decoding, the tap/swipe split
111 rule, and class statistics are in Appendix A.

112 3.1 Two evaluation protocols

113 **Exploratory protocol (Phases 1–7, historical).** Our original runs streamed the AITW training split
114 in mirror order, filtered by data mix, and took *consecutive slices*: the first n_{train} steps for training
115 and the next n_{val} for validation. This is the kind of split that looks innocuous in code review and is
116 not: episodes and, worse, repeated task goals straddle the boundary (Section 5.2). Every historical
117 number in this paper is labeled exploratory and is re-analyzed with dependence-aware statistics.

118 **Confirmatory protocol (Phase 9, frozen).** We rebuilt the data pipeline around a frozen manifest:
119 the first 12,000 `all_with_coords` steps of the pinned mirror revision, grouped by *normalized task*
120 *goal* (episodes never span goals), split goal-disjointly into train (9,511 steps), validation (1,660),
121 and an **untouched test region** (829): test groups are drawn exclusively from beyond stream index
122 5,500, past the furthest index any historical run ever loaded (5,250), so no test step was trained on,
123 evaluated on, or inspected during the exploratory phase. Evaluation sets are frozen and identical
124 for every variant and training size (250 validation / 600 test steps for `all_with_coords`; 200/400
125 for the control mix); training subsets are nested prefixes. The manifest, split hashes, 17 disclosed
126 near-duplicate goal pairs (token Jaccard ≥ 0.8), duplicate checks, and the prespecified analysis
127 (contrasts, metrics, corrections, failed-run rules) were committed to version control *before* any test-
128 set evaluation ran. We claim “decisions frozen before test access,” documented by commit history –
129 not third-party preregistration.

130 4 Methods

131 **Stage 1.** A frozen Qwen2-VL-2B encodes screenshot and instruction; mean-pooled final hidden
132 states feed a 3-layer MLP trained with class-weighted cross-entropy (full details in Appendix A).

133 4.1 Stage 2: five conditioning mechanisms

134 Stage 2 fine-tunes Qwen2-VL-2B with LoRA (rank 16 on q/k/v/o projections, $\approx 0.2\%$ of parameters)
135 to emit the coordinate string; all variants share data, optimizer ($\text{lr } 2 \times 10^{-5}$, 2 epochs, batch 1), and
136 evaluation. The only difference is how action-type information enters:

137 **A (flat baseline).** No conditioning; replicates joint decoding.

138 **B (auxiliary loss).** A linear action-type head on the pooled representation adds cross-entropy to the
139 grounding loss ($\lambda=1$); *no action information at inference*.

140 **C (hard routing).** Trained to emit “{word} (x, y)”; at inference the gold action word is forced
141 as the decode prefix.

142 **D-hook (additive embedding).** A learned table $E \in \mathbb{R}^{8 \times d}$; the row for the current action type is
143 added to *all* token embeddings via a forward hook, preserving the normal `input_ids` path (embed-
144 ding norm at convergence ≈ 0.07 at the headline setting).

145 **D-token (prepended token – the original hypothesis).** A sentinel `<|action_slot|>` token is
146 prepended to the instruction and its embedding replaced in place by $E[a]$, M-RoPE-correctly (the
147 token is in `input_ids`); norm at convergence ≈ 2.2 , a strong routable signal.

148 All variants minimize the autoregressive loss over coordinate tokens, with variant-specific condi-
149 tioning; B adds a weighted action-type cross-entropy term. The deployable end-to-end configuration
150 trains D-hook, then a Stage 1 head on frozen-backbone features, and grounds under *predicted* types;
151 gold-type results are oracle ceilings throughout (Appendix A).

152 4.2 Evaluation metrics

153 **hit@ r** : fraction of predictions within normalized L2 distance r of the target ($r \in \{0.05, 0.10, 0.25\}$;
154 headline 0.10). **Mean L2**: mean normalized distance (parse failures scored at the $\sqrt{2}$ sentinel and
155 counted as misses; parse rates reported). Statistical methodology is Section 5.3.

156 5 Three Evaluation Pitfalls, Quantified

157 5.1 Pitfall 1: M-RoPE vs. `inputs_embeds`

158 Qwen2-VL assigns 3D (temporal, height, width) position IDs to visual tokens via an `input_ids`-
159 keyed path. Our first D-token implementation injected the action embedding by passing
160 `inputs_embeds` directly, which silently bypasses that path: image-token positions were computed
161 without the action slot in the sequence, misaligning multimodal position embeddings for *every* vi-
162 sual token. The observed effect was a confident 8σ “A beats D” result – a pure artifact. Rebuilding
163 the same mechanism through the `input_ids` path (D-hook, D-token) erased the deficit. Any con-
164 ditioning study on M-RoPE-family models that touches `inputs_embeds` should verify position-ID
165 computation explicitly; the artifact is silent, large, and directionally misleading.

166 5.2 Pitfall 2: stream-slice splits leak task goals

167 Table 5 quantifies the historical protocol’s leakage by reconstructing, for every historical cell,
168 which episodes and normalized task goals each validation step shared with that cell’s training slice.
169 Episode-level overlap – the leakage one usually checks – is small (2% at the headline cell). The
170 dominant channel is *shared task goals*: 24% of headline validation steps have a goal string that also
171 appears in training, rising to 85% at $n_{\text{train}}=5000$ – AITW goals repeat across many episodes, so
172 consecutive-slice splitting almost guarantees template overlap at scale. Val-set difficulty therefore
173 varies with n_{train} (each size had a different val slice), quietly confounding the historical scaling curve.
174 The confirmatory protocol removes both channels by construction (goal-disjoint splits; Section 3.1).

175 5.3 Pitfall 3: pseudo-replicated significance

176 Our original headline significance pooled per-example distances across seeds – treating 3×250
177 (seed, example) rows as 750 independent paired observations. Predictions of the same example
178 by different seeds are strongly correlated, and examples within an episode share state; the pooled
179 bootstrap therefore overstates evidence. On synthetic nulls matching the data’s dependence structure
180 (episode and example effects shared across seeds; zero true difference), the pooled test rejects 21.2%
181 of the time at nominal 5% (1000 simulations), while the replacement below achieves 5.1%.

182 **Repaired inference.** For each contrast we (i) average each example’s paired difference over seeds,
183 (ii) cluster examples by episode, (iii) bootstrap clusters (10,000 resamples) for 95% CIs, and (iv) use
184 an episode-level sign-flip permutation test (10,000 permutations; p -values reported no smaller than
185 Monte Carlo resolution – never $p=0$). **Estimand.** This tests the *fixed trained checkpoints* on new
186 data from the eval population. It is *not* calibrated for claims about the training procedure: adding
187 run-level (seed) offsets to the synthetic null drives the episode-clustered test to 55% rejection (and
188 the pooled test to 70%). With only three seeds per variant we therefore release per-seed means for
189 every contrast (in `rerun_analysis.json` and `reanalysis.json`) and make no procedure-level
190 significance claims. The retracted pooled p -values (including a headline $p=0.0002$) should not be
191 cited.

192 6 Exploratory Results (historical splits, repaired statistics)

193 All numbers in this section come from the leaky historical splits and are therefore *exploratory*;
194 they are shown with episode-clustered statistics, and Section 7 gives the prespecified untouched-test
195 versions of every primary claim.

variant	hit@0.10	L2 ↓	parse	Δ vs. A [95% CI]	p (Holm)
A (flat)	0.219 ± 0.021	0.450	100.0%	—	—
B (aux loss)	0.230 ± 0.005	0.445	100.0%	+0.011 [-0.005, +0.029]	0.23 (0.309)
C (hard routing)	0.219 ± 0.024	0.498	91.7%	+0.000 [-0.022, +0.021]	1
D-hook (additive)	0.276 ± 0.004	0.376	100.0%	+0.057 [+0.035, +0.079]	$< 10^{-4}$ (≤ 0.0004)
D-token (prepended)	0.265 ± 0.004	0.388	100.0%	+0.046 [+0.027, +0.067]	0.0002

Table 1: **Confirmatory** main ablation on the untouched, goal-disjoint AITW test set ($n=600$ steps; $n_{\text{train}}=1200$; 3 seeds; mean $\pm$ std across seeds). Δ vs. A: episode-clustered 95% CI and permutation p (Holm-adjusted across the four primary contrasts where applicable).

196 The Stage 1 action-type classifier study (Table 7) is summarized by Figure 1 and given in full in
197 Appendix C.

198 Re-scoring the historical headline with the repaired statistics (Appendix C) leaves B–A significant
199 (+0.051 [+0.020, +0.081], $p = 0.0043$, versus the invalid pooled $p = 0.0002$), weakens D-hook
200 to marginal after correction, and renders D-token vs. B non-significant ($-0.036 [-0.077, +0.005]$,
201 $p = 0.114$) — i.e. the pooled test alone had manufactured the course-era refutation of our own
202 mechanism.

203 The exploratory per-class decomposition (Appendix C, Section C.3) attributed the advantage to
204 *click*; Section 7.1 revises that. The historical controls are also in that appendix: the AITW
205 *taps_and_swipes* mix (superseded by Section 7.2) and a Stage-2 Multimodal-Mind2Web com-
206 parison in which D-hook versus A is null (-0.006 hit@0.10 over 2 seeds), a cross-dataset null we
207 retain.

208 The exploratory low-data grid (Appendix C) shows no clean monotone trend and is confounded by
209 per-size validation sets.

210 Under clustered statistics the historical end-to-end “win” is already not significant (+0.016 [-0.017,
211 +0.052], $p = 0.424$; Appendix C).

212 A sensitivity analysis feeding the trained D-token model wrong or zeroed action embeddings (Ap-
213 pendix C, Table 8) shows the embedding is causally used at inference rather than ignored; it is a
214 sensitivity test, not evidence that conditioning caused any gain over the baseline.

215 7 Confirmatory Results (untouched test, prespecified)

216 All variants were retrained from scratch on the frozen goal-disjoint splits (identical hyperparameters,
217 seeds 42–44) and evaluated once on the untouched 600-step test set; the four primary contrasts, the
218 headline metric (hit@0.10), the clustered test, and Holm correction were committed before any test
219 evaluation ran (Section 3.1).

220 **P1 (auxiliary supervision) — null.** B vs. A on hit@0.10: +0.011 [-0.005, +0.029], permutation
221 $p = 0.23$, Holm $p = 0.309$; mean L2 $-0.005 [-0.016, +0.006]$ ($p = 0.42$). The variant the
222 exploratory analysis identified as the winner does *not* replicate. Auxiliary action-type supervision,
223 which adds no information at inference, leaves grounding statistically unchanged on data whose task
224 goals were never seen in training.

225 **P2 (inference-time conditioning) — positive, but oracle.** D-hook vs. A: +0.057 [+0.035, +0.079]
226 (Holm $p = \leq 0.0004$; mean L2 $-0.074 [-0.100, -0.050]$), positive in all three seeds (per-seed deltas
227 +0.035 to +0.085). This is the one large effect that survives the clean split. It is an *oracle* result:
228 D-hook receives the ground-truth action type at test time, which A never sees. P4 is what happens
229 when that oracle is replaced by a real classifier.

230 **P3 (the proposal’s mechanism) — reversed in our favour, and we flag it as such.** D-token vs.
231 B: +0.035 [+0.017, +0.054] (Holm $p = 0.0027$) — the prepended learned action embedding our
232 original proposal advocated significantly *outperforms* the auxiliary loss on the clean test set. The

exploratory data pointed the other way (Section 6), and the course-era version of this work reported the embedding as refuted. On the untouched test both embedding mechanisms beat B (D-token vs. B +0.035 [+0.017, +0.054]; D-hook vs. B +0.046 [+0.026, +0.066], $p = <10^{-4}$) and both beat A (D-token vs. A +0.046 [+0.027, +0.067], $p = 0.0002$), while hard routing is exactly null (C vs. A +0.000 [-0.022, +0.021], $p = 1$).

Two cautions on reading P3. First, it is an *oracle-versus-deployable* comparison, not a like-for-like mechanism ranking: D-token and D-hook read the ground-truth action type at inference while B reads nothing, so the embedding variants are being credited with information B never receives. Second, we neither prespecified nor ran a D-hook versus D-token contrast, so we make no claim that either embedding route is better than the other. What P3 does establish is that the course-era verdict — “the prepended embedding is refuted” — does not survive a clean split. We stress that this reverses in the direction that flatters our own hypothesis, which is precisely why it needed a prespecified test on data we had never touched.

P4 (deployable pipeline) — null. End-to-end predicted-type pipeline vs. A: +0.015 [-0.004, +0.034] (Holm $p = 0.309$). The reason is Stage 1 generalization, not Stage 2: the action-type classifier scores 0.795 accuracy on the frozen validation split but only 0.708 on the untouched test region, and the paired predicted-vs-oracle gap is -0.048 [-0.061, -0.035] ($p = <10^{-4}$). That gap absorbs 76% of the oracle arm’s advantage over A (0.282 ± 0.005 versus 0.219 ± 0.021 , i.e. +0.063), leaving a residual that is statistically indistinguishable from zero. **Stage-1 error leaves no detectable benefit**, directly contradicting the historical claim that “classifier error is not the bottleneck.” We note what this drop is *not*: both the frozen validation split and the untouched test region are goal-disjoint from training, so novel task goals alone cannot explain the difference between them. The two held-out sets differ in where they sit in the episode stream (the test region lies entirely beyond index 5,500), and we did not run the ablations needed to isolate that cause, so we report the drop as an observed generalization gap rather than a diagnosed one. Either way it is the binding constraint on the deployable pipeline.

7.1 Where the confirmatory gain lives: proportionally scroll, absolutely still click

Table 2 revises the exploratory mechanism story of Section C.3. On leaky data the conditioning advantage looked like *click* disambiguation. On the untouched test set the largest *relative* effect is on *scroll* (+0.098 absolute, $2.5\times$ relative) versus a smaller proportional click gain (+0.053, $1.2\times$). Because click is roughly three times as frequent in the test set, click nevertheless still supplies 61% of the aggregate gain against scroll’s 39%: what moved is where the *effect* is concentrated, not which class contributes the most additional hits. The mechanism this supports is different and simpler: swipe targets occupy a narrow, largely goal-independent region of the screen, so being told “this step is a scroll” relocates the prediction in a way the flat baseline cannot infer from the screenshot and instruction alone. B, which gets the same supervision only as a training signal, captures none of it and slightly harms scroll.

7.2 The control condition does not behave as designed

Our prespecified control assumed that on *taps_and_swipes* action type carries no spatial information, so conditioning should not help. It does help: D-hook vs. A on the control test set is +0.047 [+0.010, +0.079] ($p = 0.0327$), while B is null (+0.013 [-0.013, +0.041], $p = 0.364$) and C is significantly harmful (-0.067 [-0.100, -0.039], $p = <10^{-4}$). These control contrasts were prespecified as *secondary*, and their p -values are unadjusted; we use them to *withdraw* a scope claim rather than to establish one, which is the direction in which an unadjusted secondary result can safely be read. We report this as a **failed control**, not a second confirmation. It is coherent with Section 7.1: if the exploitable signal is “swipes land in a predictable region,” then a tap-versus-swipe mix is exactly a setting where action type *is* spatially informative, and our control was mis-designed rather than uninformative. The consequence for the paper is concrete: we cannot claim “conditioning helps if

281 and only if action type is spatially predictive,” because we no longer have a matched condition in
282 which it demonstrably is not.

283 7.3 Low data

284 The confirmatory low-data grid (Table 4) evaluates every training size against the *same* frozen test
285 set, removing the per-size-val-set confound that made Figure 4 unreadable. D-hook’s oracle advan-
286 tage is present at $n=300$ and $n=800$ and not resolved at $n=500$; B is positive only at $n=300$. There
287 is no clean monotone low-data trend in either variant, so we again decline the “low-data prior” claim
288 our proposal predicted.

289 Secondary contrasts, per-class decompositions, parse rates, and per-seed values for every run are
290 in the released `results/phase9_rerun/rerun_analysis.json`, which also carries the per-seed
291 mean for each side of every contrast. No procedure-level significance claim is made from three
292 seeds (Section 5.3).

293 7.4 What the broken evaluation actually cost

294 Because the same five-variant ablation was run under both protocols, we can separate the two defects
295 and ask what each one cost. Three states are distinguishable: (i) the original analysis (leaky split,
296 pooled bootstrap) — what a course-era version of this paper reported; (ii) the statistically repaired
297 analysis on the *same* leaky data (Section 6); and (iii) the confirmatory analysis on the untouched
298 split.

- 299 • **Which variant wins.** (i) B beats A, $p=0.0002$. (ii) B still beats A, $+0.051$ [$+0.020$,
300 $+0.081$], $p=0.0043$. (iii) *B is null* ($+0.011$ [-0.005 , $+0.029$]); the embedding mechanisms
301 are the only ones that help.
- 302 • **Whether our own hypothesis was refuted.** (i) D-token is significantly worse than B
303 (pooled $p=0.0192$), so the proposal’s prepended embedding is refuted. (ii) the gap is no
304 longer significant (-0.036 [-0.077 , $+0.005$], $p=0.114$) — the refutation was already an
305 artifact of the pooled test. (iii) *D-token significantly beats B* ($+0.035$ [$+0.017$, $+0.054$]).
- 306 • **Where the gain lives.** (i)–(ii) click disambiguation. (iii) the *relative* effect is largest on
307 *scroll* ($2.5\times$ versus $1.2\times$ for click), though click’s higher frequency still makes it 61% of
308 the absolute gain.
- 309 • **Whether the pipeline is deployable.** (i) the predicted-type pipeline was reported as beat-
310 ing A, with “classifier error is not the bottleneck” — a descriptive comparison that no test
311 ever supported. (ii) any calibrated test shows it was never significant ($+0.016$ [-0.017 ,
312 $+0.052$], $p = 0.424$). (iii) the pipeline is *null*, with Stage-1 error absorbing 76% of the
313 oracle gain.
- 314 • **Whether the control holds.** (i)–(ii) conditioning looks neutral where action type was
315 assumed uninformative. (iii) the control shows a significant effect, so it is *not a valid*
316 *control*.

317 Two lessons separate cleanly. First, the statistical repair alone was *not* sufficient: it withdrew two
318 unsupported claims (the D-token refutation and the deployable pipeline) but left the headline “auxil-
319 iary supervision works” and the click-mechanism story intact and wrong — those required the clean
320 split. Fixing inference without fixing the data would have produced a paper that was better calibrated
321 and still substantively mistaken. Second, the direction of error is not systematic: the flawed proto-
322 col inflated some effects, suppressed our own hypothesis, and mis-located the mechanism. A study
323 can be internally consistent, multi-seed, ablated, controlled, and significant at the pooled $p=0.0002$
324 we now retract, and still have every one of its mechanistic claims move under a clean split and a
325 calibrated test.

8 Limitations

Every positive effect among the four prespecified contrasts is oracle-gated: D-hook and D-token receive ground-truth action types that a deployed agent would have to predict, and when we predict them the effect disappears (P4). No headline claim in this paper should be read as a deployable improvement. One non-oracle positive does appear among the secondary results — B over A at $n_{\text{train}}=300$ (Table 4) — but it is a secondary, unadjusted contrast in a grid where the same comparison is null at both larger sizes, so we do not build a claim on it. Our control condition failed (Section 7.2), so we have no valid matched setting in which action type is demonstrably uninformative; the scope condition on when conditioning helps is therefore untested, and designing a genuine control (for instance, targets randomized within class) is future work. The scroll-versus-click mechanism reading in Section 7.1 is a post-hoc decomposition of confirmatory data, not a prespecified test.

Offline step-level hit@radius is not long-horizon task success: it scores logged single steps, without interaction, recovery, or compounding-error dynamics – the motivating failure mode is measured here only via its step-level signature. The backbone is a single 2B model; conclusions are about this scale and training recipe, not about GUI agents generally, and our baselines are small matched-compute models trained by us, not state-of-the-art systems. The confirmatory estimand is fixed checkpoints, not the training procedure (three seeds; per-seed effects reported). AITW’s type coordinates are degenerate, so “grounding” on `all_with_coords` is effectively click/scroll grounding plus a constant; the click-concentration analysis is post-hoc on exploratory data and only its prespecified aggregate versions are confirmatory. Goal strings are our dependence proxy; 17 near-duplicate train/test goal pairs (disclosed in the manifest) survive normalization, and subtler sharing (same app, similar screens) is not modeled. The Mind2Web control used 2 seeds without per-example logs. Finally, the mirror we pin is a third-party repackaging of AITW; we verified revision stability but not byte-level provenance against the original release.

9 Conclusion

On a test set whose task goals were never seen in training, the factorization thesis survives only in oracle form. Telling a 2B GUI agent the ground-truth action type improves grounding (+0.057 [+0.035, +0.079] hit@0.10). Both mechanisms that consume that oracle signal beat the auxiliary loss, including the prepended action embedding our proposal advocated (D-token–B +0.035 [+0.017, +0.054]) — a reversal of the course-era verdict that the embedding was refuted, though the comparison is oracle-versus-deployable, and we neither prespecified nor ran a D-hook versus D-token contrast, so we do not claim either embedding mechanism is better than the other. But auxiliary supervision alone is null, hard routing is null, the largest *relative* effect moves to *scroll* rather than the click disambiguation we previously reported (click still supplies 61% of the absolute gain), our control condition turns out not to be a control, and the only configuration one could actually deploy — predicting the action type and conditioning on the prediction — is statistically indistinguishable from doing nothing, because the classifier drops from 0.795 to 0.708 accuracy between the frozen validation split and the untouched test region. The honest summary is that we measured a real but oracle-gated effect and no deployable one.

The methodological result is the durable one. We ran the same five-variant ablation twice: once under a protocol with three defects (an `inputs_embeds/M-RoPE` interaction fabricating an 8σ false negative, goal-level split leakage reaching 85%, and a pooled bootstrap with 21.2% measured type-I error), and once under a repaired one. The defects did not merely exaggerate significance; they inverted which mechanism wins, which class carries the gain, whether our hypothesis was refuted, whether the pipeline was deployable, and whether the control held. None of the three is exotic — stream-sliced splits, embedding-injection shortcuts, and pooled per-seed bootstraps are ordinary practice — so we expect the same reversals are latent in other offline GUI-grounding evaluations. The frozen splits, per-example predictions, clustered-inference code, and prespecification are released so that both the object-level result and the audit can be checked and rerun.

Contributions and disclosure

Author contributions and the generative-AI usage statement are omitted from this anonymized submission and will be included in any camera-ready version. Both are present in the public preprint.

References

- [1] M. Ahn et al. Do as I can, not as I say: Grounding language in robotic affordances. In *CoRL*, 2022.
- [2] A. Brohan et al. RT-2: Vision-language-action models transfer web knowledge to robotic control. *arXiv preprint arXiv:2307.15818*, 2023.
- [3] K. Cheng et al. SeeClick: Harnessing GUI grounding for advanced visual GUI agents. *arXiv preprint arXiv:2401.10935*, 2024.
- [4] X. Deng et al. Mind2Web: Towards a generalist agent for the web. In *NeurIPS*, 2023.
- [5] B. Gou et al. Navigating the digital world as humans do: Universal visual grounding for GUI agents. *arXiv preprint arXiv:2406.01729*, 2024.
- [6] K. Q. Lin et al. ShowUI: One vision-language-action model for GUI visual agent. *arXiv preprint arXiv:2411.17465*, 2024.
- [7] Z. Lu et al. UI-R1: Enhancing action prediction of GUI agents by reinforcement learning. *arXiv preprint arXiv:2503.21620*, 2025.
- [8] T. Luo, L. Logeswaran, J. Johnson, and H. Lee. Visual test-time scaling for GUI agent grounding. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025. arXiv:2505.00684.
- [9] C. Rawles et al. Android in the wild: A large-scale dataset for android device control. In *NeurIPS*, 2023.
- [10] P. Wang et al. Qwen2-VL: Enhancing vision-language model’s perception of the world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024.
- [11] Z. Wu et al. OS-Atlas: A foundation action model for generalist GUI agents. *arXiv preprint arXiv:2410.23218*, 2024.
- [12] S. Yao et al. ReAct: Synergizing reasoning and acting in language models. *arXiv preprint arXiv:2210.03629*, 2022.
- [13] Y. Yao et al. UI-TARS: Pioneering automated GUI interaction with native agents. *arXiv preprint arXiv:2501.12326*, 2025.
- [14] Q. Zheng et al. GUI-Actor: Coordinate-free visual grounding for GUI agents. *arXiv preprint arXiv:2503.15197*, 2025.

A Method details

A.1 Sources and preprocessing

We use two datasets under a canonical 8-class action taxonomy (*click*, *double_click*, *type*, *scroll*, *drag*, *hotkey*, *wait*, *finished*).

AITW [9] (via the `cjfcst/AITW_General` HuggingFace mirror, revision pinned to `5c0dc713`) is the primary Stage-2 benchmark: five active classes (*click* 55%, *scroll* 13%, *finished* 12%, *type* 11%, *hotkey* 9%). Screenshots in the mirror are raw RGB bytes without headers; we decode them against seven observed resolutions with a portrait-aspect fallback. `DUAL_POINT` events are split into taps (*click*) versus swipes (*scroll*) by normalized touch–lift distance (threshold 0.04). Two data mixes matter below: `all_with_coords` (tap, swipe, *type* – action *type* is spatially predictive because taps target small elements while swipes are gestures) and `taps_and_swipes` (a control where knowing the *type* carries little information about *where*).

Mind2Web [4] (multimodal release) provides the Stage-1 cross-dataset comparison and a Stage-2 cross-dataset control; it is *click/type*-dominated with only two active classes at scale.

variant	click ($n=393$)	scroll ($n=136$)	type ($n=71$)
A (flat)	0.312	0.064	0.000
B (aux loss)	0.337	0.042	0.000
C (hard routing)	0.309	0.074	0.000
D-hook (additive)	0.366	0.162	0.000
D-token (prepended)	0.367	0.108	0.000

Table 2: Confirmatory per-class hit@0.10 on the untouched test set (seed-averaged; parse failures counted as misses, so denominators are the fixed class counts in the header). C, D-hook and D-token receive the ground-truth action type at inference; A and B do not. The oracle conditioning gain is proportionally largest on *scroll* (A 0.064 $\rightarrow$ D-hook 0.162, a $2.5\times$ relative improvement) and modest on *click* (0.312 $\rightarrow$ 0.366). The auxiliary loss (B) improves click slightly while *harming* scroll. *type* remains exactly 0.000 for every variant: AITW type coordinates are degenerate.

variant	hit@0.10	L2 $\downarrow$	parse	Δ vs. A [95% CI]	p
A (flat)	0.302 ± 0.009	0.268	100.0%	—	—
B (aux loss)	0.316 ± 0.018	0.264	100.0%	+0.013 [-0.013, +0.041]	0.364
C (hard routing)	0.236 ± 0.064	0.454	84.7%	-0.067 [-0.100, -0.039]	$< 10^{-4}$
D-hook (additive)	0.349 ± 0.020	0.213	100.0%	+0.047 [+0.010, +0.079]	0.0327

Table 3: Confirmatory control mix (*taps_and_swipes*) on its frozen 400-step test set. The control was designed as the falsifiable half of the thesis — conditioning should not help where action type is uninformative — and it *fails*: D-hook shows a significant gain here too (Section 7.2).

422 **Coordinates.** Targets are normalized to $[0, 1]^2$ and rendered as the string "(x, y)" on a $[0, 999]$
423 grid; training loss is masked to the coordinate tokens.

424 A.2 Stage 1: action-type classifier

425 A frozen Qwen2-VL-2B encodes screenshot+instruction; mean-pooled final hidden states (1536-
426 d) feed a 3-layer MLP trained with class-weighted cross-entropy (AdamW, lr 10^{-4} , batch 128, 8
427 epochs, best epoch by validation macro-F1). Modes: *text_only* vs. *vision_text*.

428 **Objective.** With screenshot s , instruction g , action type a , coordinate tokens $y_{1:T}$ and variant-
429 specific conditioning c :

$$\mathcal{L}_{\text{coord}} = - \sum_{t=1}^T \log p_{\theta}(y_t \mid y_{<t}, s, g, c), \quad (1)$$

430 with $c = \emptyset$ for A and B, $c = E[a]$ for D-hook/D-token, and the forced word for C; B adds
431 $\lambda \text{CE}(h_{\phi}(\bar{z}), a)$ (Eq. 1 otherwise unchanged).

432 **End-to-end pipeline.** The deployable configuration trains D-hook, then a small MLP head on
433 frozen-backbone features (Stage 1 in-job), and evaluates grounding under *predicted* types; gold-
434 type results are oracle ceilings, never deployable claims.

435 B Additional confirmatory tables

436 These tables report confirmatory results on the untouched test set and are placed here only for length;
437 they are not exploratory.

438 C Exploratory material (historical splits)

439 Everything in this appendix comes from the historical, leaky splits (Section 3.1) and is exploratory:
440 the Stage 1 classifier study, the per-class and low-data decompositions, the end-to-end and D-token

n_{train}	A	$\Delta \text{ B-A [95\% CI]}$	$\Delta \text{ D-hook-A [95\% CI]}$
$n=300$	0.138	+0.038 [+0.013, +0.064]	+0.049 [+0.025, +0.073]
$n=500$	0.219	-0.004 [-0.027, +0.018]	+0.016 [-0.008, +0.039]
$n=800$	0.256	+0.003 [-0.014, +0.020]	+0.036 [+0.015, +0.057]

Table 4: Confirmatory low-data grid on the *same* frozen test set for every training size (nested training prefixes), removing the per-size-val-set confound of Figure 4.

historical cell	val steps	sharing a train episode	sharing a train goal
all_wc, $n=300$	250	4 (2%)	14 (6%)
all_wc, $n=500$	250	4 (2%)	47 (19%)
all_wc, $n=800$	250	5 (2%)	49 (20%)
all_wc, $n=1200$	250	4 (2%)	61 (24%)
all_wc, $n=2500$	250	10 (4%)	133 (53%)
all_wc, $n=5000$	250	2 (1%)	213 (85%)
taps_swipes, $n=1000$	200	0 (0%)	48 (24%)

Table 5: Leakage in the historical (exploratory) splits: fraction of validation steps sharing an episode or a normalized task goal with that cell’s training slice. Goal sharing, not episode straddling, is the dominant channel.

441 sensitivity analyses, and three figures. It is superseded for inferential purposes by the confirmatory
442 results in the main text.

443 C.1 Main ablation under clustered statistics

444 Table 6 re-scores the historical headline with the repaired statistics (per-seed values in Figure 3).
445 Three observations. First, the auxiliary-loss effect survives: B-A on hit@0.10 is +0.051 [+0.020,
446 +0.081] ($p = 0.0043$), where the invalid pooled analysis had claimed $p = 0.0002$ – on *this* data the
447 conclusion stands, with the evidence overstated by an order of magnitude. (Section 7 shows it does
448 not survive a clean split at all: fixing the statistics was not sufficient.) Second, D-hook weakens
449 to marginal after multiplicity correction. Third, the comparison our proposal cared about most –
450 D-token vs. B – is *no longer significant* ($-0.036 [-0.077, +0.005]$, $p = 0.114$): on the historical data
451 we cannot even claim the prepended embedding is significantly *worse* than the auxiliary loss, only
452 that it shows no sign of being better.

453 C.2 Stage 1: vision is essential exactly where instructions under-specify the action

454 Table 7 and Figure 1: on Mind2Web, instructions name the action (“click the submit button”), so
455 vision adds +0.009 macro-F1 over text; on AITW, instructions state goals (“open a new Chrome
456 private window”), text-only collapses toward the majority class, and vision adds +0.414. A TF-IDF
457 ceiling of 0.168 on AITW confirms text is uninformative there regardless of classifier strength.

variant	hit@0.10	L2 ↓	$\Delta \text{ vs. A [95\% CI]}$	p
A (flat)	0.255 ± 0.026	0.392	—	—
B (aux loss)	0.305 ± 0.015	0.362	+0.051 [+0.020, +0.081]	0.0043
C (hard routing)	0.279 ± 0.039	0.412	+0.024 [-0.016, +0.063]	0.278
D-hook (additive)	0.300 ± 0.035	0.365	+0.045 [+0.011, +0.085]	0.0227
D-token (prepended)	0.269 ± 0.031	0.375	+0.015 [-0.018, +0.047]	0.437

Table 6: Exploratory main ablation (AITW all_with_coords, $n_{\text{train}}=1200$, 3 seeds, historical leaky splits). Episode-clustered $\Delta \text{ vs. A}$ with 95% cluster-bootstrap CIs and sign-flip permutation p . Under the repaired analysis, B survives (Holm-adjusted $p = 0.017$ across the four vs.-A contrasts); D-hook is marginal after correction (Holm $p = 0.068$); C and D-token are not significant.

Variant	Mind2Web macro-F1	AITW macro-F1
text_only	0.597 $\pm$ 0.011	0.141 $\pm$ 0.032
vision_text	0.605 $\pm$ 0.016	0.555 $\pm$ 0.036
vision_zeroed	0.469 $\pm$ 0.000	0.110 $\pm$ 0.051
Vision delta	+0.009	+0.414
TF-IDF + LogReg	0.622	0.168
Majority floor	0.472	0.140

Table 7: Stage 1 action-type classification (3 seeds, mean $\pm$ std; exploratory splits). On Mind2Web the instruction text alone contains the action type; on AITW it does not, and vision contributes +0.414 macro-F1. `vision_zeroed` collapsing to the majority floor is the sanity check that features, not priors, drive the result.

Per-class decomposition (AITW all_with_coords, n=1200)




Figure 2: Exploratory hit@0.10 by action class (all_with_coords). The conditioning advantage concentrates in *click*; *type* is degenerate on AITW (all variants score exactly 0.000, mean L2 at the $\approx \sqrt{2}$ corner-to-corner distance).

Unlike the Stage-2 experiments, these Stage-1 runs are *not* affected by the stream-slicing defect of Section 5.2: they train on each dataset’s own train split and evaluate on its held-out test split (AITW) or test_task split (Mind2Web). We did not, however, audit goal overlap *between* those official splits, so we treat these numbers as descriptive context rather than as a contamination-checked result. The Stage-1 accuracy that matters for the deployable claim is measured on our own frozen splits in Section 7, where it is substantially lower on unseen goals.

This asymmetry decides where Stage-2 conditioning could matter at all.

C.3 Per-class decomposition: the gain is click disambiguation

Breaking hit@0.10 down by class (Figure 2; deltas below are recomputed on matched seeds 42–44 with parse failures counted as misses, which the original Phase-7 aggregation did neither of): B gains +0.079 on click (68% of examples) at a small scroll cost (-0.014); C gains +0.075 on click but destroys scroll (-0.145) by forcing a decode prefix that derails the gesture class; D-hook alone improves both (click +0.057, scroll +0.036). The *type* class scores exactly 0.000 across all variants with mean normalized L2 at or slightly above $\sqrt{2}$, the corner-to-corner distance (a few predictions land outside the unit square): AITW *type* events carry degenerate touch coordinates and are ungroundable by any model. This corrects our earlier “*type*→text-field localization” interpretation: the real signal is **click-vs-scroll/type disambiguation** for the precise-target class, and the “*type*” rows merely add a constant drag that cancels in deltas.

contrast	Δ hit@0.10 [95% CI]	Δ hit@0.25 [95% CI]	p
wrong – gold	-0.192 [-0.233, -0.149]	-0.329 [-0.372, -0.284]	$< 10^{-4}$
zero – gold	-0.093 [-0.128, -0.059]	-0.112 [-0.145, -0.079]	$< 10^{-4}$

Table 8: Exploratory D-token sensitivity (episode-clustered, 3 seeds): evaluating the same trained model with wrong or zeroed action embeddings severely degrades grounding. This establishes the embedding is *used* at inference – a sensitivity/oracle analysis, not evidence that conditioning caused any gain over A.

476 C.4 Control conditions

477 On `taps_and_swipes` (action type spatially uninformative), the historical 3-seed means show no
478 conditioned variant meaningfully above A (A 0.390 ± 0.008 , D-hook 0.392 ± 0.035 – a difference
479 well inside one seed-std – with B 0.368 ± 0.009 and C 0.325 ± 0.032 at or below A on hit@0.10);
480 per-example logs were not retained for A on this mix, so the clustered version of this control appears
481 only in the confirmatory rerun (Table 3). On Multimodal-Mind2Web (click/type-dominated), D-
482 hook vs. A was -0.006 hit@0.10 / $+0.018$ hit@0.25 over 2 seeds – a null we retain deliberately.
483 On the historical data both of these looked like confirmations that conditioning does not help where
484 action type was *assumed* to carry no spatial information. Section 7.2 shows that assumption is wrong
485 for the AITW control, which therefore never tested what we believed it tested; the Mind2Web null
486 stands on its own as a cross-dataset result.

487 C.5 Low-data behavior: no clean curve

488 Our proposal predicted the advantage is a low-data prior. The historical grid ($n \in \{300, 500, 800\}$,
489 3 seeds) does *not* show a clean monotone: B–A on hit@0.10 is $+0.067$ at $n=300$ but $+0.011$ at
490 $n=500$ and -0.040 at $n=800$; D-hook is more stable ($+0.069$, $+0.001$, $+0.048$) and leads hit@0.25
491 and mean L2 at every low-data size. The two single-seed high-end cells disagree with each other:
492 at $n=5000$ the deltas shrink and partially reverse (B–A hit@0.25 goes negative), while at $n=2500$
493 D-hook posts the *largest* conditioning deltas anywhere in the study ($+0.076$ hit@0.10, $+0.204$
494 hit@0.25) – single-seed noise on a differently-leaky val slice, cutting in the direction *favorable*
495 to conditioning. We report all of this as: *inconclusive on the low-data-prior hypothesis – noisy,*
496 *non-monotone, and confounded by per-size val sets* (Figure 4); the prespecified low-data claim is
497 evaluated on the frozen test set in Section 7. An earlier draft of this work called the low-data
498 prediction “confirmed”; that claim did not survive this audit.

499 C.6 The e2e pipeline claim also weakens

500 The historical end-to-end result (predicted-type pipeline 0.271 vs. A 0.255 hit@0.10) was previ-
501 ously reported as a win. Under clustered statistics the paired difference is $+0.016$ [-0.017 , $+0.052$]
502 ($p = 0.424$) – not significant; the oracle-type version is marginal ($+0.037$ [$+0.005$, $+0.076$]). The
503 deployable-pipeline question is therefore decided by the confirmatory run (Section 7, P4), not the
504 historical data.

505 C.7 D-token causal-use sensitivity

506 Feeding the trained D-token model a *wrong* action id collapses grounding, and zeroing the embed-
507 ding hurts substantially (Table 8). What this establishes is narrow: the prepended embedding is
508 causally *used* at inference rather than ignored, so any verdict against it cannot be explained away
509 as an inert conditioning path. Whether it is in fact the worse way to exploit action-type supervision
510 was the course-era conclusion, and Section 7 (P3) reverses it on clean data. Figure 5 shows one
511 example where correct conditioning concentrates more attention mass on the screenshot; across the
512 three examples we probed the pattern is not uniform, so we present it strictly as illustration, not
513 evidence.

Stage 2 ablation (AITW all_with_coords): $B \approx D\text{-hook} > C > A$




Figure 3: Exploratory ablation on all_with_coords (historical splits; per-seed values). Shown for continuity with the historical record; Table 6 carries the calibrated statistics and Table 1 the confirmatory test-set versions.

Data-scaling: action-type conditioning advantage vs training size
(AITW all_with_coords; 300-1200 are 3-seed means, 2500/5000 single seed)




Figure 4: Exploratory conditioning advantage ($\Delta \text{hit@0.10}$ vs. A) across n_{train} . Caveat from Section 5.2: each historical size used a *different* val slice with different goal leakage (up to 85%), so this curve is directional at best; the confirmatory grid (Table 4) fixes one frozen test set for all sizes.



Figure 5: Image-region attention for a held-out *click* example under three action-type conditionings (qualitative illustration; screenshot attention mass 0.204 under *click* vs. 0.132/0.110 under *type/scroll*).